

# Approximation Algorithms: Load Balancing

# Approximations

- The word *approximation* is used for several flavors of algorithms
- Approximating optimization problems
  - An algorithm is a  $c$ -approximation if its cost is  $c \cdot \text{OPT}$ , where OPT is the optimum cost; ( $c < 1$  or a maximization problem,  $c > 1$  for min)
- Approximations in an online setting
  - An online algorithm is a  $c$ -approximation/ $c$ -competitive, if it has cost  $c \cdot \text{OPT}$ , where OPT is the cost of an offline algorithm (that knows the entire input ahead of time); ( $c < 1$  or a max problem,  $c > 1$  for min)
- The word approximate is also used to indicate that the algorithm is permitted to make one-sided errors (such false positive)
  - Bloom filter
  - Approximate hitter algorithm (Problem 5(b) on assignment)

# Challenges: Approximation Algorithms

- Approximating problems that are NP hard
  - Main challenge is showing that the algorithm performs close to optimal when the optimal solution is not known/NP hard
  - Usually done by lower (upper) bounding the cost of the optimal solution for minimization (maximization) problems
- Approximation for online algorithms
  - High benchmark. Comparison against an optimal that knows the entire future, while the algorithm does not even know the next element
  - Sometimes dealt with using “resource augmentation”—allowing the algorithm some flexibility compared to the optimal

# Online: Ski Rental Problem

- Assume that you are taking ski lessons
- After each lesson you decide (depending on how much you liked it and how cold you are) whether to continue to ski or to quit entirely
- Question: rent or buy?
- Cost of renting \$1 (say)
- Cost of buying \$ $B$
- **Offline strategy.** If you knew in advance how many times you would ski, say  $t$  times, what is the best strategy?
  - If  $t \geq B$  times, then buy, else rent
  - In other words, optimal offline cost is  $\min\{t, B\}$



# Online: Ski Rental Problem

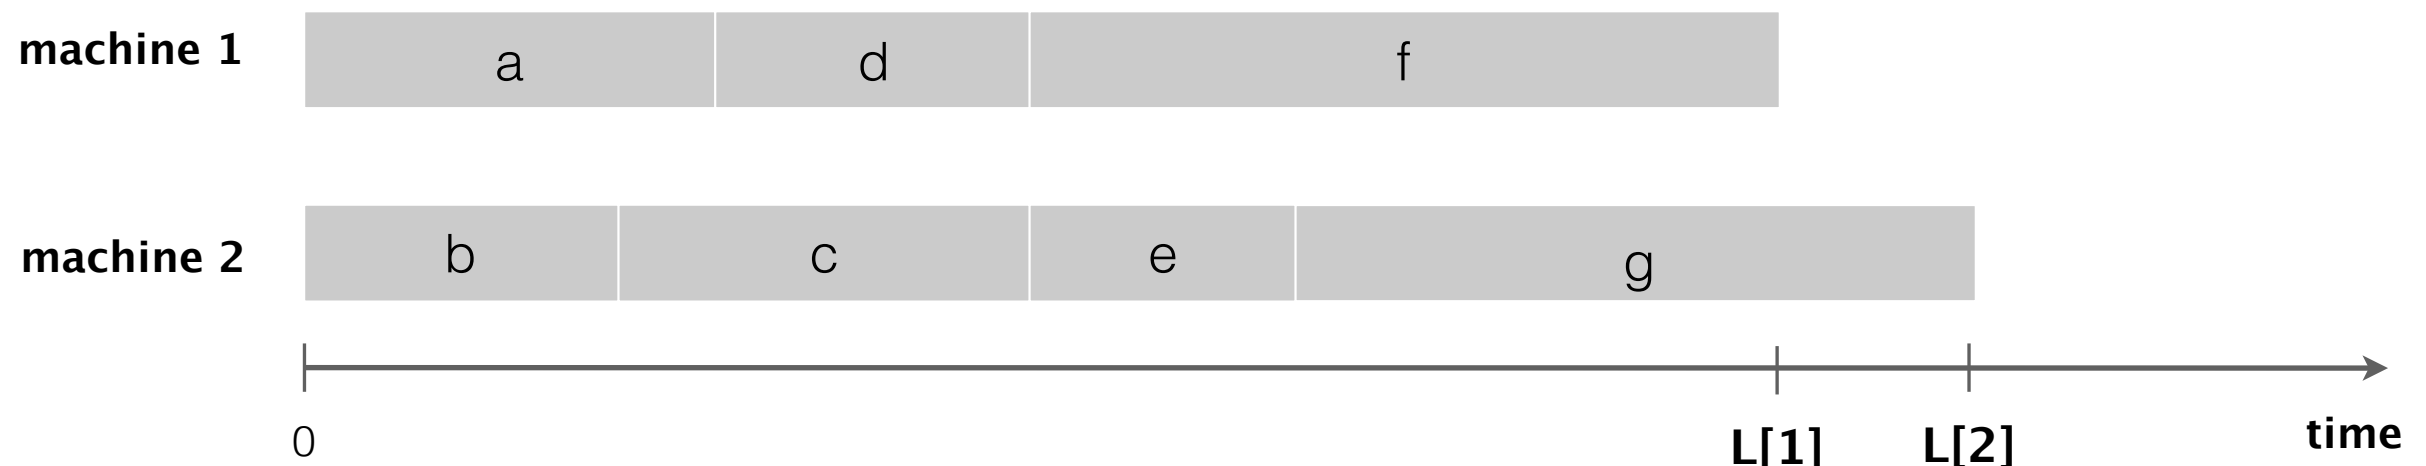
- **Online strategy.** We need to figure out a decision point, a number  $k$  such you buy skis on the  $k$ th visit (renting before then)
- **Claim.** If we set  $k = B$  (the cost of buying skis), we are guaranteed to never pay more than twice of the best offline optimal strategy
- That is, buying on the  $B$ th ski visit is 2-competitive
- Even if you quit right after the  $B$ th visit,  $t \geq B$
- Offline cost is  $\min\{t, B\} = B$
- Online strategy's cost?
  - $(k - 1) \cdot 1 + B = (B - 1) + B = 2B - 1$
- Competitive/Approximation ratio?
  - $2B - 1/B = 2 - 1/B \leq 2$



# Load Balancing

- **Input.**  $m$  identical machines;  $n$  jobs, and processing times  $t_1, \dots, t_m$ , where job  $j$  has processing time  $t_j$  (on any machine)
- Job  $j$  must run contiguously on one machine
- A machine can process at most one job at a time.
- Let  $S[i]$  be the subset of jobs assigned to machine  $i$ .

- The **load of machine  $i$**  is  $L[i] = \sum_{j \in S[i]} t_j$  (total processing time).



# Load Balancing

- The **makespan** of an algorithm is the maximum load on any machine  
$$L = \max_i L[i]$$
- **Load balancing Problem.** Assign jobs to machines so as to minimize makespan.
- **Claim.** Load balancing is NP hard even with  $m = 2$  machines
- **Proof.** Reduction from PARTITION problem.
- We will design an approximation algorithm for this problem
- [Greedy returns!] Consider the following greedy strategy:
  - Fix some order on the jobs
  - Assign job  $j$  to machine  $i$  whose load is smallest so far

# Load Balancing: Greedy

**LIST-SCHEDULING** ( $m, n, t_1, t_2, \dots, t_n$ )

**FOR**  $i = 1$  **TO**  $m$

$L[i] \leftarrow 0.$   $\leftarrow$  load on machine  $i$

$S[i] \leftarrow \emptyset.$   $\leftarrow$  jobs assigned to machine  $i$

**FOR**  $j = 1$  **TO**  $n$

$i \leftarrow \operatorname{argmin}_k L[k].$   $\leftarrow$  machine  $i$  has smallest load

$S[i] \leftarrow S[i] \cup \{j\}.$   $\leftarrow$  assign job  $j$  to machine  $i$

$L[i] \leftarrow L[i] + t_j.$   $\leftarrow$  update load of machine  $i$

**RETURN**  $S[1], S[2], \dots, S[m].$

- Running time?
  - $O(n \log m)$  using a priority queue for loads  $L[k]$



# Load Balancing: Greedy Analysis

- **Claim.** Greedy algorithm is a 2-approximation.
- To show this, we need to show greedy solution never more than a factor two worse than the optimal
- **Challenge.** We don't know the optimal solution. In fact, finding the optimal is NP hard.
- Technique used in approximation algorithm (minimization problem)
  - Lower bound the cost of optimal solution
  - A good enough lower bound can help show that our algorithm cannot be too much worse than the optimal
- In our problem, what are some lower bounds on the makespan of even an optimal algorithm?

# Load Balancing: Greedy Analysis

- Let OPT be the optimal makespan.

- **Lemma.**  $\text{OPT} \geq \max_j t_j$ .

- **Proof.** Some machine must process the most time-consuming job.
- Any other lower bounds?

- **Lemma.**  $\text{OPT} \geq \frac{1}{m} \sum_j t_j$


- **Proof.**

- The total processing time is  $\sum_j t_j$

- Some machine must do a  $1/m$  fraction of the total work.

# Greedy is a 2-Approximation

- **Proof.** Consider load  $L[i]$  of bottleneck machine  $i$  ← machine that ends up with highest load
  - Let  $j$  be the last scheduled job on machine  $i$
  - When job  $j$  was assigned to machine  $i$ ,  $i$  must have had the smallest load
  - That is,  $L[i] - t_j \leq L[k] \quad \forall 1 \leq k \leq m$



# Greedy is a 2-Approximation

- **Proof.** Consider load  $L(i)$  of bottleneck machine  $i$ 
  - Let  $j$  be the last scheduled job on machine  $i$
  - When job  $j$  was assigned to machine  $i$ ,  $i$  must have had the smallest load
  - That is,  $L[i] - t_j \leq L[k] \quad \forall 1 \leq k \leq m$
  - Summing over all  $k$  and dividing by  $m$

$$\begin{aligned} L[i] - t_j &\leq \frac{1}{m} \sum_k L[k] \\ &\quad \frac{1}{m} \sum_k t_k \\ &\leq \text{OPT} \end{aligned}$$

# Greedy is a 2-Approximation

- **Proof.**


- Consider load  $L(i)$  of bottleneck machine  $i$

- $$L[i] - t_j \leq \frac{1}{m} \sum_k L[k]$$

$$\frac{1}{m} \sum_k t_k \leq \text{OPT}$$

- We know that  $t_j \leq \text{OPT}$

- Thus,  $L = L[i] \leq \text{OPT} + t_j \leq 2\text{OPT}$  ■



# Greedy is a 2-Approximation

- Is our analysis tight?
- Close to it.
- Consider  $m(m - 1)$  jobs of length 1 + 1 job of length  $m$
- How would greedy schedule these jobs?
  - Greedy will evenly divide the first  $m(m - 1)$  jobs among  $m$  machines, will place the final long job on any one machine
  - Makespan:  $m - 1 + m = 2m - 1$
- How would optimal schedule it?
  - Give the long job to one machine, the rest split the other small jobs with a makespan  $m$
- Ratio:  $(2m - 1)/m \approx 2$

# Greedy is Online

- Notice that our greedy algorithm is an online algorithm
- Assigns jobs to machines in the order they arrive
  - Does not depend on future jobs
- Online approximation algorithms are very useful as often the entire input is not known ahead of time
- In online settings, it may be impossible to compute an optimum solution in polynomial time, even when the offline problem is polynomial time solvable
- Can we do better, if we assume all jobs are available at start time?
- **Offline.** Slight modification of greedy gets better approximation!

# Improving on Online Greedy

- Worst case of our greedy algorithm: spreading jobs out evenly when a giant job at the end screwed things up
- What can we do to avoid this?
  - Idea: deal with larger jobs first
  - Small jobs can only hurt so much
- Turns out this improves our approximation factor
- **Longest-processing-time (LPT) first.** Sort  $n$  jobs in decreasing order of processing times; then run the greedy algorithm on them
- **Claim.** LPT has a makespan at most  $1.5 \cdot \text{OPT}$
- **Observation.** If we have fewer than  $m$  jobs, then the greedy solution is clearly optimal (as it puts each job on its own machine)




# LPT-first is a 1.5-Approximation

- **Lemma.** LPT-first has a makespan at most  $1.5 \cdot \text{OPT}$
- **Observation.**
  - If we have fewer than  $m$  jobs, then the greedy solution is clearly optimal (as it puts each job on its own machine)
- **Claim.** If more than  $m$  jobs then,  $\text{OPT} \geq 2 \cdot t_{m+1}$
- **Proof.** Consider the first  $m + 1$  jobs in sorted order.
  - They each take at least  $t_{m+1}$  time
  - $m + 1$  jobs and  $m$  machines, there must be a machine with at least two jobs
  - Thus the optimal makespan  $\text{OPT} \geq 2 \cdot t_{m+1}$

# LPT-first is a 1.5-Approximation

- **Lemma.** LPT-first has a makespan at most  $1.5 \cdot \text{OPT}$
- **Proof.** Similar to our original proof. Consider the machine  $M_i$  that has the maximum load
- If  $M_i$  has a single job, then our algorithm is optimal
- Suppose  $M_i$  has at least two jobs and let  $t_j$  be the last job assigned to the machine, note that  $j \geq m + 1$  (why?)

- Thus,  $t_j \leq t_{m+1} \leq \frac{1}{2}\text{OPT}$




# LPT-first is a 1.5-Approximation

- **Lemma.** LPT-first has a makespan at most  $1.5 \cdot \text{OPT}$
- **Proof.** Similar to our original proof. Consider the machine  $M_i$  that has the maximum load
- If  $M_i$  has a single job, then our algorithm is optimal
- Suppose  $M_i$  has at least two jobs and let  $t_j$  be the last job assigned to the machine, note that  $j \geq m + 1$  (why?)

- Thus,  $t_j \leq t_{m+1} \leq \frac{1}{2}\text{OPT}$

- $T_i - t_j \leq \text{OPT}$

- $T_i \leq \frac{3}{2}\text{OPT}$  ■



# Is our 1.5-Approximation tight?

- **Question.** Is our  $3/2$ -approximation analysis tight?
  - Turns out, no
- **Theorem [Graham 1969].** LPT-first is a  $4/3$ -approximation.
  - Proof via a more sophisticated analysis of the same algorithm
- **Question.** Is the  $4/3$ -approximation analysis tight?
  - Pretty much.
- Example
  - $m$  machines,  $n = 2m + 1$  jobs
  - 2 jobs each of length  $m, m + 1, \dots, 2m - 1$  + one job of length  $m$
  - Approximation ratio  $= (4m - 1)/3m \approx 4/3$

# Formal Definition

- Consider an arbitrary optimization problem
- Let  $\text{OPT}(X)$ : the cost of the optimal solution on a given input  $X$
- Let  $A(X)$ : the cost of algorithm  $A$  on the same input  $X$
- $A$  is a  $\alpha(n)$ -approximation iff
$$\frac{\text{OPT}(X)}{A(X)} \leq \alpha(n) \quad \text{and} \quad \frac{A(X)}{\text{OPT}(X)} \leq \alpha(n)$$
for all input  $X$  of size  $n$
- **Maximization problem**: second inequality is trivial, first matters
- **Minimization problem**: first inequality is trivial, second matters
- Goal. Find a useful function of the input to upper and lower on the cost of  $\text{OPT}$  and  $A$ , e.g.  $\text{OPT}(X) \geq f(X)/2$  and  $A(X) \leq 4f(X)$  means  $A$  is a 8-approximation

# Acknowledgments

- Some of the material in these slides are taken from
  - Kleinberg Tardos Slides by Kevin Wayne (<https://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/04GreedyAlgorithmsI.pdf>)
  - Jeff Erickson's Algorithms Book (<http://jeffe.cs.illinois.edu/teaching/algorithms/book/Algorithms-JeffE.pdf>)
  - Lecture slides: <https://web.stanford.edu/class/archive/cs/cs161/cs161.1138/>