

Advanced Algorithms

Module 2

Parallel Computing



Why Parallel Computing?

- Save time, resources, memory, ...
- Who is using it?
 - Academia
 - Industry
 - Government
 - Individuals?
- Two practical motivations:
 - Application requirements
 - Architectural concerns.
- Why now?
 - Most computers including laptops are multi-core!
 - Need to therefore study how to use parallel computers.

Conventional Wisdom in Computer Architecture

- Power Wall + Memory Wall + ILP Wall = Brick Wall
- Old CW: Uniprocessor performance 2X / 1.5 yrs
- New CW: Uniprocessor performance only 2X / 5 yrs?

The Academic Interest

- Algorithmics and complexity
 - How to design parallel algorithms?
 - What are good theoretical models for parallel computing?
 - How to analyze parallel algorithms?
 - Can every sequential algorithm be parallelized?
 - What are some complexity classes wrt parallel computing?

The Academic Interest

- Systems and Programming
 - How to write parallel programs?
 - What are some tools and environments.
 - How to convert algorithms to efficient implementations.
 - What are the differences to sequential programming?
 - What are the performance measures?
 - Can sequential programs be automatically converted to parallel programs?


The Academic Interest

- Architectures
 - What are standard architectural designs?
 - What new issues are raised due to multiple cores?
 - Downstream concerns
 - Does a programmer have to worry about this?
 - How to support the systems software as architecture changes?

The Course Coverage

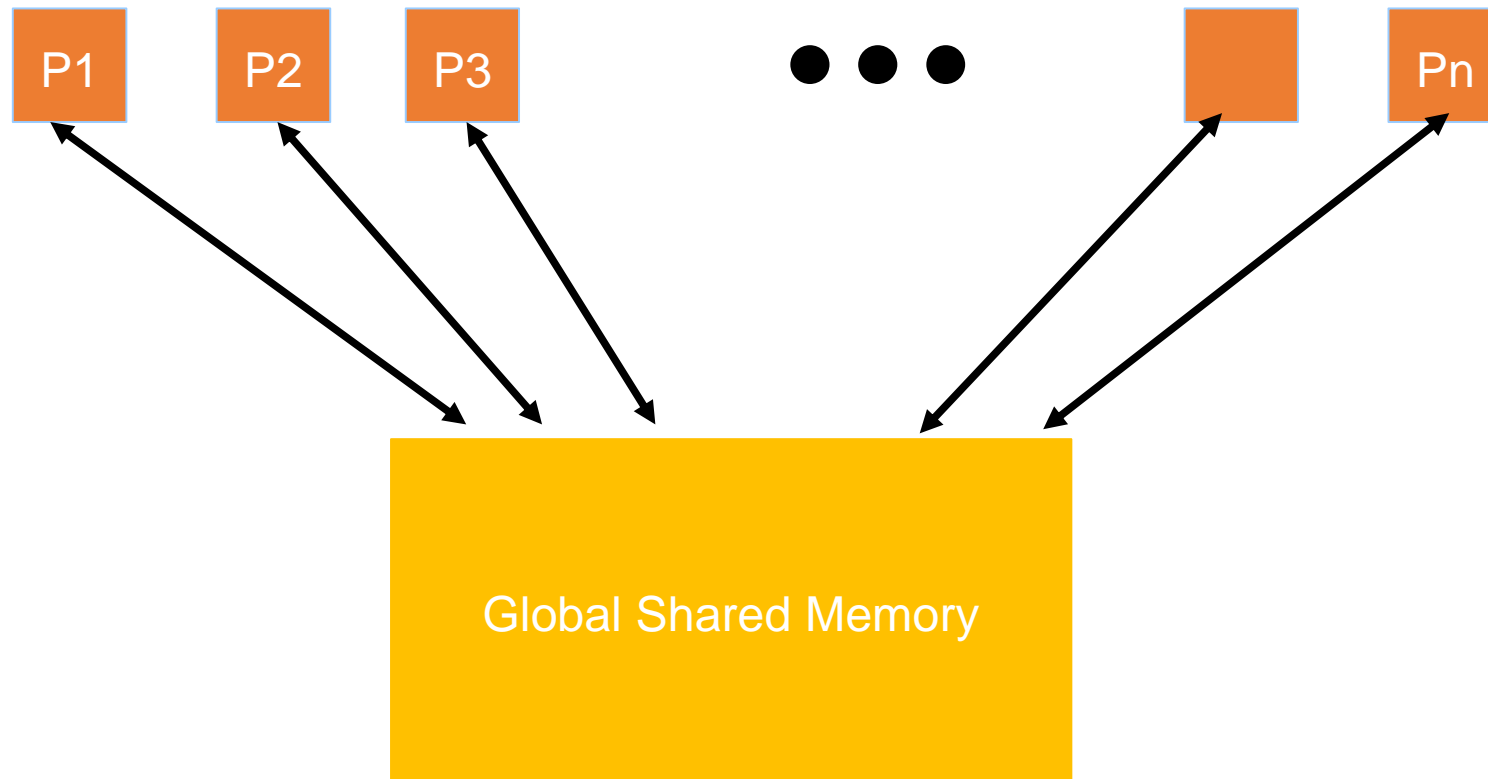
- Focus on algorithms and complexity
- Models for parallel algorithms
- Algorithm design methodologies with application
 - Semi-numerical
 - Lists
 - Trees and graphs
- Complexity, characterization, and connection to sequential complexity classes.

The PRAM Model



- The von Neumann model.

The PRAM Model



- An extension of the von Neumann model.

The PRAM Model

- A set of n identical processors
- A common access shared memory
- Synchronous time steps
- Access to the shared memory costs the same as a unit of computation.
- Different models to provide semantics for concurrent access to the shared memory
 - EREW, CREW, CRCW(Common, Arbitrary, Priority, ...)

The Semantics

- In all cases, it is the programmer to ensure that his program meets the required semantics.
- EREW : Exclusive Read, Exclusive Write
 - No scope for memory contention.
 - Usually the weakest model, and hence algorithm design is tough.
- CREW : Concurrent Read, Exclusive Write
 - Allow processors to read simultaneously from the same memory location at the same instant.
 - Can be made practically feasible with additional hardware

The Semantics

- **CRCW : Concurrent Read, Concurrent Write**
 - Allow processors to read/write simultaneously from/to the same memory location at the same instant.
 - Requires further specification of semantics for concurrent write. Popular variants include
 - **COMMON** : Concurrent write is allowed so long as the all the values being attempted are equal. Example: Consider finding the Boolean OR of n bits.
 - **ARBITRARY** : In case of a concurrent write, it is guaranteed that some processor succeeds and its write takes effect.
 - **PRIORITY** : Assumes that processors have numbers that can be used to decide which write succeeds.

PRAM Model – Advantages and Drawbacks

Advantages

- A simple model for algorithm design
- Hides architectural details for the designer.
- A good starting point

Disadvantages


- Ignores architectural features such as:
 - memory bandwidth,
 - communication cost and latency,
 - scheduling, ...
- Hardware may be difficult to realize

Example 1 – Matrix Multiplication

- One of the fundamental parallel processing tasks.
- Applications to several important problems in linear algebra, signal processing and optimization.
- Several techniques that work in parallel also.

Example I – Matrix Multiplication

- Recall that in $C = A \times B$, $C[i,j] = \sum A[i,k].B[k,j]$.
- Consider the following recursive approach:
 - Works well in practice.



$$C_{00} = A_{00} \cdot B_{00} + A_{01} \cdot B_{10}$$


$$C_{01} = A_{00} \cdot B_{01} + A_{01} \cdot B_{11}$$

$$C_{10} = A_{10} \cdot B_{00} + A_{11} \cdot B_{10}$$

$$C_{11} = A_{10} \cdot B_{01} + A_{11} \cdot B_{11}$$

Example I – Matrix Multiplication

- Other approaches include Cannon's algorithm



- Can overlap computation with communication.
- Works well when the number of processors is more.

Example 2 – New Parallel Algorithm

Listing 1:

$S(1) = A(1)$

for $i = 2$ to n do

$S(i) = S(i-1) \circ A(i)$

- **Prefix Computations:** Given an array A of n elements and an associative operation \circ , compute $A(1) \circ A(2) \circ \dots \circ A(i)$ for each i .
- A very simple sequential algorithm exists for this problem.
- Many computations can be expressed in terms of prefix computations.

Parallel Prefix Computation

- The sequential algorithm in Listing 1 is not efficient in parallel.
 - In particular, has to wait for the output of $S(i)$ to compute the output $S(i+1)$.
- Need a new algorithm approach.
 - Balanced Binary Tree


Balanced Binary Tree

- An **algorithm design approach** for parallel algorithms
- Many problems can be solved with this design technique.
- Easily amenable to parallelization and analysis.

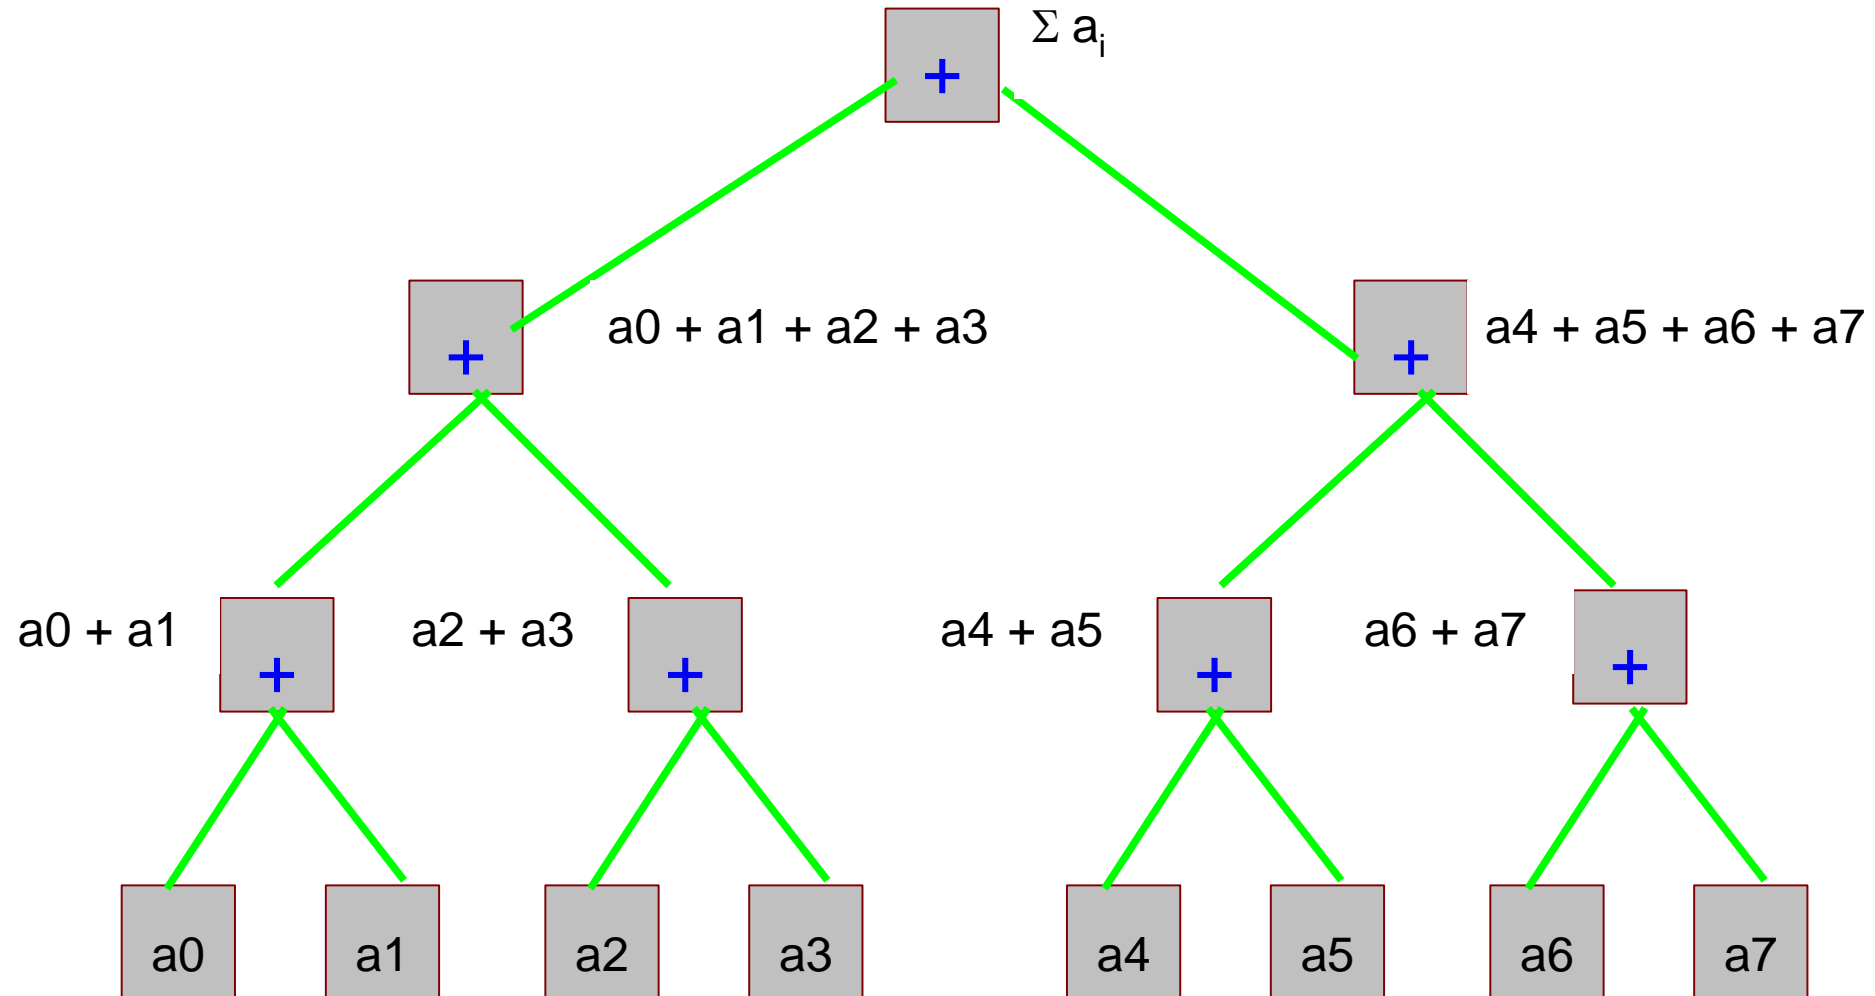
Balanced Binary Tree

- A complete binary tree with processors at each internal node.
- Input is at the leaf nodes
- Define operations to be executed at the internal nodes.
 - Inputs for this operation at a node are the values at the children of this node.
- Computation as a tree traversal from leaf to root.

Balanced Binary Tree – Prefix Sums



Balanced Binary Tree – Sum



Balanced Binary Tree – Sum

- The above approach called as an ``upward traversal"
 - Data flow from the children to the root.
 - Helpful in other situations also such as computing the max, expression evaluation.
- Analogously, can define a downward traversal
 - Data flows from root to leaf
 - Helps in settings such as element broadcast

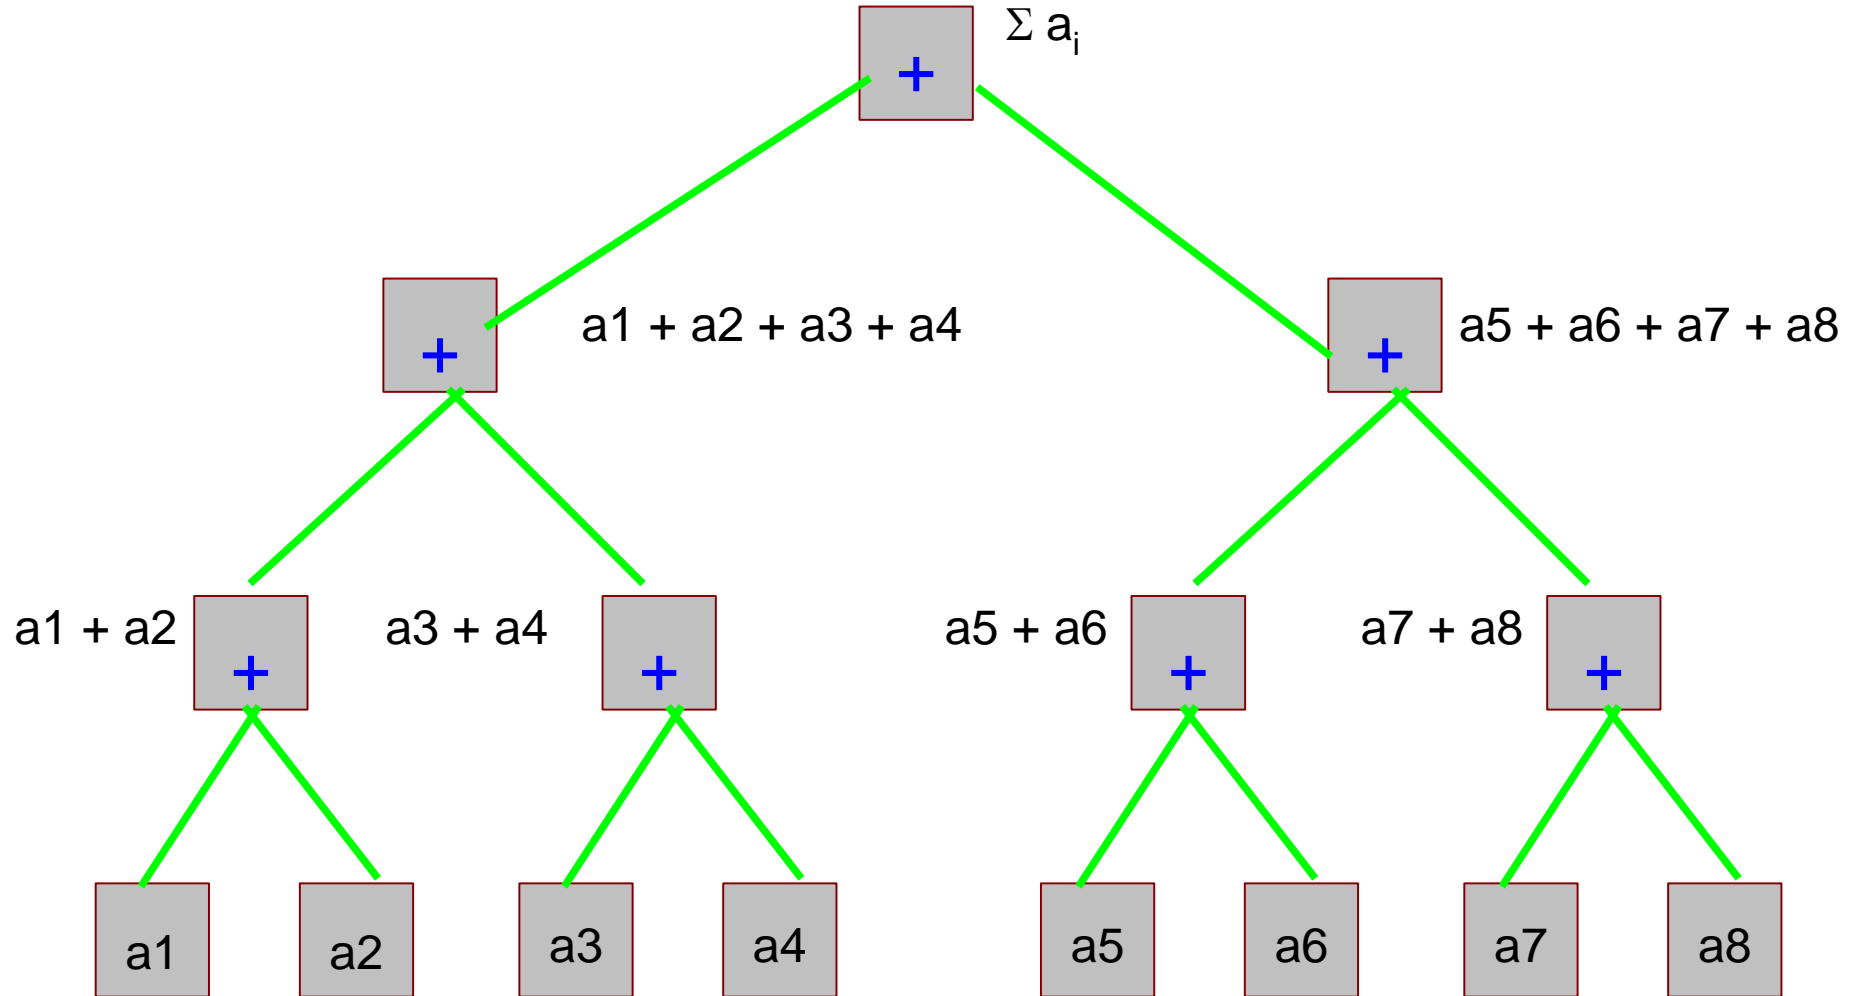
Balanced Binary Tree

- Can use a combination of both upward and downward traversal.
- Prefix computation requires that.
- Illustration in the next slide.

A : 4 2 -1 3 0 5

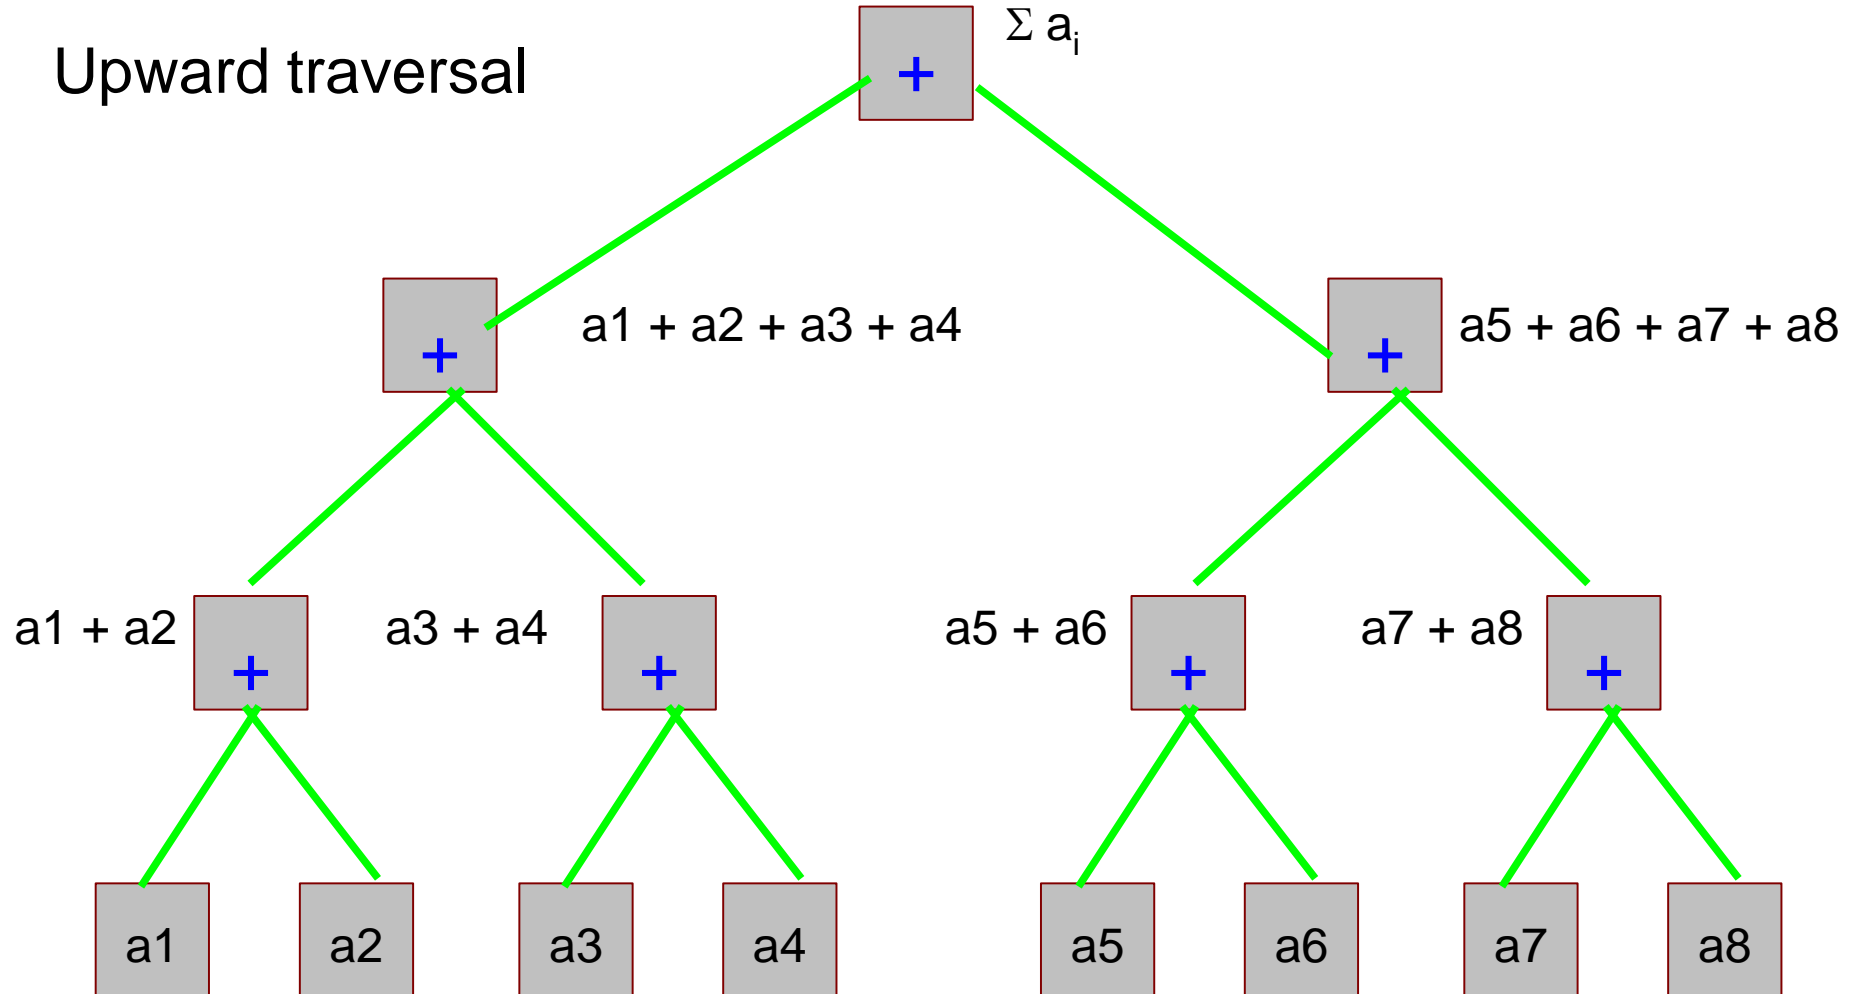
PrefixSum(A) : 4 6 5 8 8 13

Balanced Binary Tree – Sum



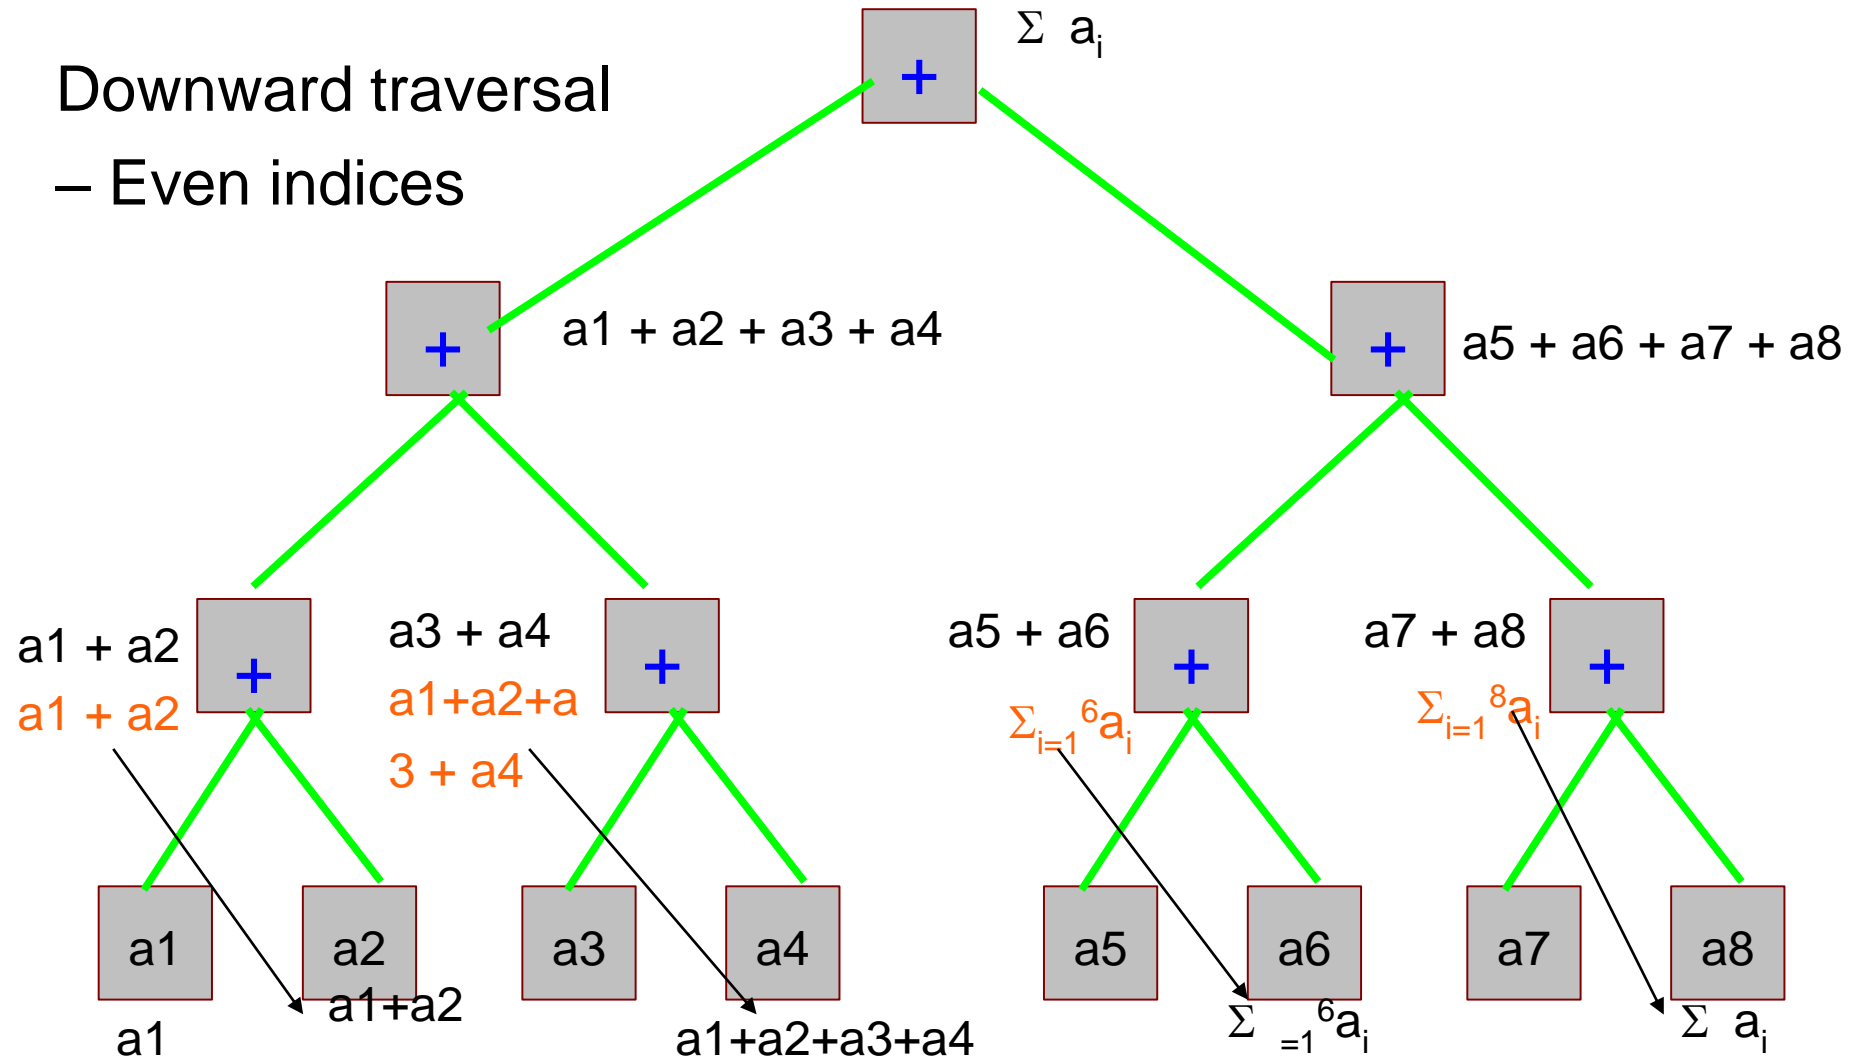
Balanced Binary Tree – Prefix Sum

Upward traversal




Balanced Binary Tree – Prefix Sum

Downward traversal
– Even indices



Balanced Binary Tree – Prefix Sum

Downward traversal
– Odd indices



Balanced Binary Tree – Prefix Sums

- Two traversals of a complete binary tree.
- The tree is only a visual aid.
 - Map processors to locations in the tree
 - Perform equivalent computations.
 - Algorithm designed in the PRAM model.
 - Works in logarithmic time, and optimal number of operations.

//upward traversal

1. for $i = 1$ to $n/2$ do in parallel

$$b_i = a_{2i-2} \circ a_{2i}$$

2. Recursively compute the prefix sums of $B = (b_1, b_2, \dots, b_{n/2})$ and store them in $C = (c_1, c_2, \dots, c_{n/2})$

//downward traversal

3. for $i = 1$ to n do in parallel

$$i \text{ is even} : s_i = c_i$$

$$i = 1 : s_1 = c_1$$

$$i \text{ is odd} : s_i = c_{(i-1)/2} \circ a_i$$

Analysis of Parallel Algorithms

- To analyze parallel algorithms, we rely on asymptotics and recurrences.
- Each operation costs 1 unit, only sequential time needs to be counted. We assume **as many processors as can be used** are available.
- In the prefix sum example, let $T(n)$ be the time in parallel for an input of size n .
 - Step 1 can use $n/2$ processors in parallel each taking 1 unit of time.
 - Step 2 is a recursive call and takes $T(n/2)$ time.
 - Step 3 uses n processors each taking 1 unit of time.

Analysis of Parallel Algorithms

- The recurrence relation is:
 - $T(n) = T(n/2) + O(1)$
 - Can ignore effects due to constant factors, such as the difference in the number of processors between steps 1 and 3.
- The solution to the above recurrence is $T(n) = O(\log n)$.
- Another parameter of interest in parallel algorithms is the work done.
- Can be stated as the sum of the works done by each of the processors.


Analysis of Parallel Algorithms

- The work done by the prefix algorithm can be expressed by the recurrence
 - $W(n) = W(n/2) + O(n)$.
 - The $O(n)$ accounts for the work in the first and the third steps.
 - Solution: $W(n) = O(n)$.
- Work done can indicate if the algorithm is doing about the same amount of operations as the best known sequential algorithm.
- Such a parallel algorithm is called an **optimal algorithm**.

Other Design Paradigms

- Partitioning
 - Similar to divide and conquer
 - But **no need** to combine solutions
 - Can treat problems independently and solve in parallel.
 - Example: Parallel merging, searching.

Merging in Parallel by Partitioning



- Two sorted arrays A and B to be merged into C.
- Let A be a sorted array. Let $\text{Rank}(x, A)$ be the number of elements smaller than x in A.
- Claim: $\text{Rank}(x, C) = \text{Rank}(x, A) + \text{Rank}(x, B)$
- For x in A, $\text{Rank}(x, A)$ is immediately available. To find $\text{Rank}(x, B)$ can use binary search in parallel.

Quick Example

A = [8 10 12 24]

B = [15 17 27 32]

Element	8	10	12	24	15	17	27	32
Rank in A	0	1	2	3	3	3	4	4
Rank in B	0	0	0	2	0	1	2	3
Rank in C	0	1	2	5	3	4	6	7

C = [8 10 12 15 17 24 27 32]

Merging in Parallel by Partitioning

- Time for each binary search is $O(\log n)$
- Total time for merging = $O(\log n)$, the total work is $O(n \log n)$.
 - Not work optimal as compared to the best possible sequential time complexity of $O(n)$.
- Can **reduce** the total work to $O(n)$.
 - Induce equal-sized partitions in the arrays
 - Rank one element, say the first element, from each partition
 - Use these ranks to find the ranks of the other elements, sequentially.


An Improved Optimal Algorithm

- **General technique**
 - Solve a smaller problem in parallel
 - Extend the solution to the entire problem.
- For the first step, the problem size to be solved is guided by the factor of non-optimality of an existing parallel algorithm.


An Improved Parallel Algorithm

- Our simple parallel algorithm is away from work optimality by a factor of $O(\log n)$.
- So, we should solve a problem of size $O(n/\log n)$.
- For this purpose, we pick every $\log n^{\text{th}}$ element of A , and similarly in B .
- Use the simple parallel algorithm on these elements of A and B .
 - Binary search however in the entire A and B .

An Improved Parallel Algorithm



- Let $A_1, A_2, \dots, A_{n/\log n}$ be the elements of A ranked in B .
- These ranks induce partitions in B .
 - Define $[B_{r(i)} \dots B_{r(i+1)}]$ as the portion of B so that $[A(i) \dots A(i+1)]$ have ranks in.
- Can therefore merge $[A(i) \dots A(i+1)]$ with $[B_{r(i)} \dots B_{r(i+1)}]$ sequentially.



An Improved Parallel Algorithm

- Such sequential merges can happen in parallel, at each index of $A[i]$.
- Time taken for the sequential merge is $O(\log n + B_{r(i+1)} - B_{r(i)})$.
- Time:
 - Binary search: $O(\log n)$, with $n/\log n$ processors.
 - Sequential merge: $O(\log n)$, subject to certain conditions. There are also $n/\log n$ such merges in parallel.
- Work:
 - There are $n/\log n$ binary searches in parallel. Work = $O(n)$.
 - For the sequential merges too, work = $O(n)$.