

Complexity and Advanced Algorithms

Spring 2021

Approximation Algorithms – A Brief Introduction



Approximation Algorithms

- Suppose a problem is known to be NP-Complete.
- No hope to solve in polynomial time unless $P = NP$.
- But, several practical problems fall in this category.
- Need some solution to this issue.
- This is where approximation algorithms help.

Approximation Algorithms

- For a problem P , let A be an approximation algorithm.
- Suppose that the problem P is a minimization problem.
- Then, the performance of Algorithm A for P is measured as its approximation ratio defined as follows.
- For an instance I of P , let $OPT(I)$ denote the best possible solution.
- Let $A(I)$ denote the solution produced by the algorithm A .
- The approximation ratio of algorithm A is $\max_I |A(I)|/|OPT(I)|$.
- Notice that the ratio is always at least 1.

Approximation Algorithms

- For a problem P , let A be an approximation algorithm.
- Suppose that the problem P is a maximization problem.
- Then, the performance of Algorithm A for P is measured as its approximation ratio defined as follows.
- For an instance I of P , let $OPT(I)$ denote the best possible solution.
- Let $A(I)$ denote the solution produced by the algorithm A .
- The approximation ratio of algorithm A is $\max_I |OPT(I)|/|A(I)|$.
- Notice that the ratio is always at least 1.

$$OPT(I) \geq A(I)$$

Approximation Algorithms

- We have seen an example of this definition earlier.
- In the context of MAXSAT.
- Today, we will study two more problems and approximation algorithms for them.

max $f(I)$

s.t. ,

best possible $I : f(\text{OPT}) = 20$

for any alg. $f(A(I)) \leq 20$

for any input, if $\frac{| \text{OPT}(I) |}{| A(I) |} \leq \frac{3}{2}$

$$| A(I) | \geq \left(\frac{2}{3} \right) | \text{OPT}(I) |$$

3/4 for MAXSAT

Load Balancing

- Consider m machines M_1, M_2, \dots, M_m , that are identical in all respects.
 - Like the m cores of your multicore computer.
- We have n jobs, J_1, J_2, \dots, J_n to be processed and any job can be processed by any machine.
- Our goal is to minimize the time spent by any machine.

Load Balancing

- We define the following quantities.
- Let $A(i)$ be the jobs assigned to machine M_i .
- Job j has a time requirement t_j , for $j = 1$ to n .
- The **makespan** of machine M_i is $T_i := \sum_{j \in A(i)} t_j$.
- The makespan of an assignment $T = \max_i T_i$.
- The goal of the problem is to find an assignment that **minimizes** the makespan.

Load Balancing

- One of the popular algorithms is to use a greedy technique.
- We can assume that all the jobs are given apriori.
 - A bit unlike the real world setting where user jobs are fired any time.
- Assign the next job to the machine that is presently least loaded.
 - Note that all jobs are given at the beginning.
 - Assignment is done before any machine starts its processing.

Load Balancing

- GreedyAssign(Jobs J, m)
- begin
 - Set $A(i)$ = empty for all i between 1 to m .
 - Set $T_i = 0$ for all i between 1 to m
 - For $j = 1$ to n do
 - Find an index i such that machine M_i has minimum T_i
 - $A(i) = A(i) \cup \{j\}$
 - $T_i = T_i + t_j$
 - Endfor
- End

Load Balancing

- Illustrate how this algorithm works on the following set of jobs and four machines.

- Runtimes of jobs = 2, 3, 2, 2, 4, 1, 2, 1.
 some machine makespan ≥ 4
max makespan ≥ 4

- Find the best possible makespan and the makespan produced by the greedy algorithm.

- For $j = 1$ to n do
 - Find an index i such that machine M_i has minimum T_i
 - $A(i) = A(i) \cup \{j\}$
 - $T_i = T_i + t_j$
- Endfor

greedy : 6
best : 5

Load Balancing

- We start with two observations that help us prove the approximation ratio of the greedy algorithm.
- Let T^* denote the best possible makespan.
- **Observation 1:** $T^* \geq (1/m) \sum_{j=1}^n t_j$.
- Comes from the fact that there is a total work of $\sum_{j=1}^n t_j$, and some machine will have to work for at least a $1/m$ fraction of the total.
- **Observation 2:** $T^* \geq \max_j t_j$.
- The longest job will be on some machine which will work for at least that much time.

Load Balancing

- Let us now turn our attention to the greedy algorithm.
- Consider the machine that has the largest makespan according to the assignment produced by the greedy algorithm.
- Let this be the machine M_i , and its makespan be T_i .
- Let t_j be the last job assigned to M_i .
- Why did we assign t_j to M_i ?
 - As M_i has the least load just before assigning t_j to M_i .

Load Balancing

- The load of M_i before assigning T_j is $T_i - t_j$.
- At this time, every other machine has a load at least $T_i - t_j$.
 - Plus any other jobs assigned to them later on.
- So, the sum of the time for all jobs is lower bounded as follows.
- $\sum_{j=1}^n t_j \geq m(T_i - t_j)$, or
- $T_i - t_j \leq (1/m) \sum_{j=1}^n t_j \leq T^*$.

Load Balancing

- Further, notice that $t_j \leq T^*$.
- Use Observation 2.

$$t_j \leq \max_k t_k \leq T^*$$

\uparrow
Obs(2)

- Therefore, $T_i = \underbrace{T_i - t_j}_{\text{}} + \underbrace{t_j}_{\text{}} = \underbrace{(T_i - t_j) + t_j}_{\text{}} \leq T^* + T^* = 2T^*.$

- Hence, $T_i \leq 2T^*.$

$$\frac{|A(I)|}{|\text{OPT}(I)|} = \frac{T_i}{T^*} \leq \frac{2T^*}{T^*} \leq 2$$

Load Balancing

- What would be one way to improve the solution produced by the greedy algorithm?
- Sort the jobs in descending order according to their runtime and assign them just as earlier.
- Let us call this as SortedGreedyAssignment.
- We will now analyze this approach.

Load Balancing

- Convince yourself that also SortedGreedyAssignment does not produce the best possible makespan.

Load Balancing

- The proof is not much different from the earlier proof.
- We will obtain a better bound for t_j , the last job assigned to the machine M_i that eventually has the largest makespan, T_i .
- Consider the case where there are fewer than m jobs.
- Then, the best possible assignment is to give each job to a different machine.
 - Our algorithm also does the same.
 - So, we indeed produce the best possible makespan.

Load Balancing

- Let us consider the case that n is more than m .
- More jobs than machines.
- In any assignment, the job t_{m+1} in sorted order, is given to some machine that already has one additional job that is at least as long as t_{m+1} .
- So, $T^* \geq 2t_{m+1}$.

Load Balancing

- Further, note that for machine M that has the largest makespan in our algorithm, if t_j is the last job assigned to M_i , then $t_j \leq t_{m+1}$.
- From the previous, $t_{m+1} \leq 1/2 T^*$.
- Now, $T_i = T_i - t_j + t_j \leq T^* + 1/2 T^* = 3/2 T^*$.