


Other Design Paradigms

- Partitioning
 - Similar to divide and conquer
 - But **no need** to combine solutions
 - Can treat problems independently and solve in parallel.
 - Example: Parallel merging, searching.

Merging in Parallel by Partitioning



- Two sorted arrays A and B to be merged into C.
- Let A be a sorted array. Let $\text{Rank}(x, A)$ be the number of elements smaller than x in A.
- Claim: $\text{Rank}(x, C) = \text{Rank}(x, A) + \text{Rank}(x, B)$
- For x in A, $\text{Rank}(x, A)$ is immediately available. To find $\text{Rank}(x, B)$ can use binary search in parallel.

Quick Example

A = [8 10 12 24]

B = [15 17 27 32]

Element	8	10	12	24	15	17	27	32
Rank in A	0	1	2	3	3	3	4	4
Rank in B	0	0	0	2	0	1	2	3
Rank in C	0	1	2	5	3	4	6	7

C = [8 10 12 15 17 24 27 32]

Merging in Parallel by Partitioning

- Time for each binary search is $O(\log n)$
- Total time for merging = $O(\log n)$, the total work is $O(n \log n)$.
 - Not work optimal as compared to the best possible sequential time complexity of $O(n)$.
- Can **reduce** the total work to $O(n)$.
 - Induce equal-sized partitions in the arrays
 - Rank one element, say the first element, from each partition
 - Use these ranks to find the ranks of the other elements, sequentially.

An Improved Optimal Algorithm


- **General technique**
 - Solve a smaller problem in parallel
 - Extend the solution to the entire problem.
- For the first step, the problem size to be solved is guided by the factor of non-optimality of an existing parallel algorithm.

An Improved Parallel Algorithm


- Our simple parallel algorithm is away from work optimality by a factor of $O(\log n)$.
- So, we should solve a problem of size $O(n/\log n)$.
- For this purpose, we pick every $\log n^{\text{th}}$ element of A , and similarly in B .
- Use the simple parallel algorithm on these elements of A and B .
 - Binary search however in the entire A and B .



An Improved Parallel Algorithm

$$\frac{n}{\log n}$$



- Let $A_1, A_2, \dots, A_{n/\log n}$ be the elements of A ranked in B .
- These ranks induce partitions in B .
 - Define $[B_{r(i)} \dots B_{r(i+1)}]$ as the portion of B so that $[A(i) \dots A(i+1)]$ have ranks in.
- Can therefore merge $[A(i) \dots A(i+1)]$ with $[B_{r(i)} \dots B_{r(i+1)}]$ sequentially.






An Improved Parallel Algorithm


- Such sequential merges can happen in parallel, at each index of $A[i]$.
- Time taken for the sequential merge is $O(\log n + B_{r(i+1)} - B_{r(i)})$.
- Time:
 - Binary search: $O(\log n)$, with $n/\log n$ processors.
 - Sequential merge: $O(\log n)$, subject to certain conditions. There are also $n/\log n$ such merges in parallel.
- Work:
 - There are $n/\log n$ binary searches in parallel. Work = $O(n)$.
 - For the sequential merges too, work = $O(n)$.

An Improved Parallel Algorithm



- What if $[B_{r(i)} \dots B_{r(i+1)}]$ has a size of more than $\log n$?
- The situation can be addressed
 - Pick equally spaced, no more than $\log n$, spaced items in $[B_{r(i)} \dots B_{r(i+1)}]$.
 - Rank these in $[A_i \dots A_{i+1}]$.

An Improved Parallel Algorithm



- What if $[B_{r(i)} \dots B_{r(i+1)}]$ has a size of more than $\log n$?
- The situation can be addressed
 - Pick equally spaced, no more than $\log n$, spaced items in $[B_{r(i)} \dots B_{r(i+1)}]$.
 - Rank these in $[A_i \dots A_{i+1}]$.

Final Result

- Can merge two sorted arrays of size n in time $O(\log n)$ with work $O(n)$.
 - Need CREW model, for binary searches.
- Can improve further, we will see later.
- The technique to achieve optimality is a general technique, with several applications. We will see more applications of this later.


A Further Improvement

- Where is the scope for improvement?
- Each binary search takes $O(\log n)$ time, and we also have $O(n/\log n)$ subproblems each of size $O(\log n)$.
- To get further improvements, we should look at both aspects.
- Can we search faster? Parallel?

A Further Improvement

- Parallel search first.
- Consider a sorted array A of n element and we want to search for an element x .
- Given p processors, we can always search at positions (indices) $1, n/p, 2n/p, \dots, n$.
- Record the result of each comparison as a 1 or 0 with 1 for position i indicating that $A[i] < x$ and 0 indicating that $A[i] \geq x$.
- The sequence of p results will have :
 - Either all 1's
 - Either all 0's
 - A shift from 1's to 0's

x : element to be searched $W(n) = W(\frac{n}{p}) + \textcircled{p}$
 $= O(\log_p^n)$




Time:

$$T(n) = T(\frac{n}{p}) + O(1)$$

$$= O(\log_p^n)$$

$$= O(\frac{\log n}{\log p})$$



A Further Improvement

- Parallel search first.
- Consider a sorted array A of n element and we want to search for an element x .
- Given p processors, we can always search at positions (indices) $1, n/p, 2n/p, \dots, n$.
- Record the result of each comparison as a 1 or 0 with 1 for position i indicating that $A[i] < x$ and 0 indicating that $A[i] \geq x$.
- The sequence of p results will have :
 - Either all 1's : x is not in A
 - Either all 0's : x is not in A
 - A shift from 1's to 0's : x is likely in the n/p segment corresponding to the shift from 1 to 0.

Search in Parallel

- We can identify the next step depending on the three cases.
 - Either all 1's : x is not in A
 - Either all 0's : x is not in A
 - A shift from 1's to 0's : x is likely in the n/p segment corresponding to the shift from 1 to 0.
 - Therefore, search recursively in the corresponding segment of size n/p while still using p processors.
- The recurrence relation for the time taken is
 - $T(n) = T(n/p) + O(1)$, for a solution of $T(n) = O(\log_p n)$.

Search in Parallel

- Consider typical values of p .
- For $p = O(1)$, no change in time taken asymptotically.
- For $p = O(\log n)$, the time taken is $O(\log n / \log \log n)$.
- For $p = O(n^{1/2})$, the time taken is $O(\log n / \log n^{1/2}) = O(1)$!
 - Of course, looks like wasteful from a work point of view.
 - Let us see what it is good for!

From Parallel Search to Merge

$0 < \epsilon < 1$ $p = n^\epsilon, T(n) = O(1/\epsilon), W(n) = O(n^\epsilon)$

- Recall our idea to arrive at an optimal algorithm to merge two sorted arrays A and B.
- We rank a few elements of A in B to partition B into sub-arrays.
- Let us consider ranking $n^{1/2}$ elements of A in B.
- We have n processors, so each search can use $n^{1/2}$ processors!
- Each search now finishes in $O(1)$ time.

$$\begin{array}{lcl}
 T(n) = O(\log_p n) & W(n) = O(p \log_p n) & \\
 p = \sqrt{n} : T(n) = O(1) & W(n) = O(\sqrt{n}) & \\
 p = O(1) : T(n) = O(\log n) & W(n) = O(\log n) & \parallel
 \end{array}$$