
Advanced Algorithms

Spring 2021

Lecture 1

Agenda for Today

- Welcome back to a new semester !!
 - We are hopefully seeing the last embers of the pandemic.
 - Nevertheless, we have to prepare for a possibly new normal in many ways.
 - We will continue to teach and learn in the online mode for Spring 2021 too.
 - Stay safe, but curious!

Agenda for Today

- Welcome back to a new semester !!
 - We are hopefully seeing the last embers of the pandemic.
 - Nevertheless, we have to prepare for a possibly new normal in many ways.
 - We will continue to teach and learn in the online mode for Spring 2021 too.
 - Stay safe, but curious!
- My details
 - Kishore Kothapalli, Professor, IIIT Hyderabad
 - Email: kkishore@iiit.ac.in (the best way to reach me)
 - Research interests span parallel computing and distributed algorithms.

Agenda for Today

- Rest of Today's Class
 - Syllabus
 - Policies
 - Expectations
 - Actual lecture

Syllabus

- Roughly, a three module course
 - Module 1: Randomized Algorithms
 - Module 2: Parallel and Distributed Algorithms
 - Module 3: Advanced Topics
 - Big data and Sampling
 - Algorithm engineering
 - Any other

Syllabus

- Roughly, a three part course
 - Randomized Algorithms
 - Chernoff bounds and Randomized Routing
 - Perfect Hashing
 - Graph algorithms: MIS, Spanners
 - Randomized Rounding
 - Approximate counting
 - Parallel and Distributed Algorithms

Syllabus

- Roughly, a three part course
 - Randomized Algorithms
 - Parallel and Distributed Algorithms
 - Flynn's Taxonomy and Models of Computations
 - Basic parallel algorithms: Search/Sort/Scan/Merge
 - Tree and Graph Algorithms
 - Lower Bounds

Policies

- Grading (Tentative)
 - Homeworks: 30% (We will have some lateness policy here)
 - In-class Quizzes : 25% (We will have some redundancy here)
 - End Exam: 15%
 - Quizzes 1 and 2 : 30%
 - Exceptional Performance: 5% extra
- Any submission that is graded and evaluated **should not** be copied from any source.
- Copied submissions will get zero for the first instance and negative for repeat offences.

We have two Teaching Assistants Support Staff named so far:

Sayantana Jana, [sayantan.jana@research.iiit.ac.in](mailto:sayantana.jana@research.iiit.ac.in)

Athreya Chandramouli, athreya.chandramouli@research.iiit.ac.in

Policies

- Textbooks: Do not own these books just for the class!
 - Randomized Algorithms, Motwani and Raghavan
 - Introduction to Parallel Algorithms, J. JaJa
 - Other material to be posted on the course website
- Most welcome to write to me if you have any questions.

Expectations

- Utilize class time effectively.
 - Starts with all of you settling by the class time.
 - Ask any question you may have. No question is small to ask.
 - Do not show up late.

Agenda for Today

- On to the actual lecture....randomization in computing
- We will see how randomization can help in designing and analyzing algorithms.
- Starting with a very simple example...

A Simple Example

- Let us recall the partition procedure and quick sort.
- We assume that the elements of the set are all distinct.



Algorithm RandQuickSort(S)

Choose a pivot element x_i u.a.r from S

Split the set S into two subsets $S_1 = \{x_j | x_j < x_i\}$
and $S_2 = \{x_j | x_j > x_i\}$ by comparing each x_j with x_i

Recurse on sets S_1 and S_2

Output the sorted set S_1 , x_i , and then sorted S_2 .

end Algorithm

A Simple Example

Algorithm RandQuickSort(S)

Choose a pivot element x_i u.a.r from S

Split the set S into two subsets $S_1 = \{x_j | x_j < x_i\}$
and $S_2 = \{x_j | x_j > x_i\}$ by comparing each x_j with x_i

Recurse on sets S_1 and S_2

Output the sorted set S_1 , x_i and then sorted S_2 .

end Algorithm

- Let $T(n)$ be the time taken by the procedure.
- What can we say about $T(n)$?

A Simple Example

Algorithm RandQuickSort(S)

Choose a pivot element x_i u.a.r from S

Split the set S into two subsets $S_1 = \{x_j | x_j < x_i\}$ and $S_2 = \{x_j | x_j > x_i\}$

Recurse on sets S1 and S2

Output the sorted set S_1 , x_i and then sorted S_2 .

end Algorithm

- What can we say about $T(n)$?
- The maximum value of $T(n)$ occurs when the pivot element x_i is the largest/smallest element of the remaining set during **each** recursive call of the algorithm.
- In this case, $T(n) = n + (n - 1) + \dots + 1 = O(n^2)$.
- This value of $T(n)$ is reached with a very low probability of $1/n \cdot 1/(n-1) \cdot \dots \cdot 1/2 \cdot 1 = 1/n!$.
- Also, the best case occurs when **every** pivot element splits the applicable set into two equal sized subsets and then $T(n) = O(n \ln n)$.

A Simple Example

Algorithm RandQuickSort(S)

Choose a pivot element x_i u.a.r from S

Split the set S into two subsets $S_1 = \{x_j | x_j < x_i\}$ and $S_2 = \{x_j | x_j > x_i\}$

Recurse on sets S1 and S2

Output the sorted set S_1 , x_i and then sorted S_2 .

end Algorithm

- This implies that $T(n)$ has a distribution between $O(n \ln n)$ and $O(n^2)$.
- Now we derive the expected value of $T(n)$.

A Simple Example

- This implies that $T(n)$ has a distribution between $O(n \ln n)$ and $O(n^2)$.
- Now we derive the expected value of $T(n)$.
- Note that if the i^{th} smallest element is chosen as the pivot element then S_1 and S_2 will be of sizes $i - 1$ and $n - i - 1$ respectively.
- And this choice has a probability of $1/n$.
- Hence, the recurrence relation for $T(n)$ is:
 - $T(n) = n + T(X) + T(n - 1 - X)$
- In the above, X is a random variable indicating the size of S_1 .

A Simple Example

- Note that if the i th smallest element is chosen as the pivot element then S_1 and S_2 will be of sizes $i - 1$ and $n - i - 1$ respectively.
- Hence, the recurrence relation for $T(n)$ is:
- $T(n) = n + T(X) + T(n - 1 - X)$
- In the above, X is a random variable indicating the size of S_1 .
- Further, note that $\Pr[X = i] = 1/n = \Pr[n - 1 - X = i]$ as $\Pr[X = i] = 1/n$.
- The last part is true since the choice of the pivot is uniform.

A Simple Example

- Hence, the recurrence relation for $T(n)$ is:
- $T(n) = n + T(X) + T(n - 1 - X)$
- Further, note that $\Pr[X = i] = 1/n = \Pr[n - 1 - X = i]$ as $\Pr[X = i] = 1/n$.
- Taking expectations on both sides,
- $E[T(n)] = n + \frac{1}{n} \sum_{i=1}^{n-1} E[T(i)] + \frac{1}{n} \sum_{i=1}^{n-1} E[T(i)]$.
- Use the fact that for a random variable Y with its support partitioned into sets A_1, A_2, \dots, A_n , we have that $E[Y] = \sum_i \Pr(A_i) \cdot E[Y \mid A_i]$.
- Let $f(i) = E[T(i)]$.

A Simple Example

- Taking expectations on both sides of,
- $E[T(n)] = n + \frac{1}{n} \sum_{i=1}^{n-1} E[T(i)] + \frac{1}{n} \sum_{i=1}^{n-1} E[T(i)]$.
- Let $f(i) = E[T(i)]$.
- We can simplify the expression as $f(n) = n + \frac{2}{n} \sum_i f(i)$.
- Further simplification results in $nf(n) = n^2 + 2(f(1) + f(2) + \dots + f(n-1))$.

A Simple Example

- Further simplification results in
- $nf(n) = n^2 + 2(f(1) + f(2) + \dots + f(n - 1))$.
- Write the above by replacing n with $n - 1$ to get

$$(n - 1) f(n-1) = (n-1)^2 + 2(f(1) + f(2) + \dots + f(n - 2))$$

- Subtract the two equations to get:
 $nf(n) = n^2 + 2(f(1) + f(2) + \dots + f(n - 1))$.

A Simple Example

- Further simplification results in
- $nf(n) = n^2 + 2(f(1) + f(2) + \dots + f(n - 1))$.
- Write the above by replacing n with $n - 1$ to get
- $(n - 1) f(n-1) = (n-1)^2 + 2(f(1) + f(2) + \dots + f(n - 2))$.
- Subtract the two equations to get:
$$nf(n) = n^2 + 2(f(1) + f(2) + \dots + f(n - 1))$$

or $f(n) = (n+1)/n f(n - 1) + (2n - 1)/n$.
- We prove by induction that $f(n) \leq 2n \ln n$.

A Simple Example

- $f(n) = (n+1)/n f(n-1) + (2n-1)/n$.
- We prove by induction that $f(n) \leq 2n \ln n$.
- Check the base case for $n = 1$.
- Let the result hold for all values of n up to $n-1$.
- Induction step:

Let the claim hold for all values up to $n-1$. Then,

$$\begin{aligned} f(n) &= \frac{n+1}{n} f(n-1) + \frac{2n-1}{n} \\ &\leq \frac{n+1}{n} 2(n-1) \ln(n-1) + \frac{2n-1}{n} \text{ by induction hypothesis} \\ &= \frac{2(n^2-1)}{n} \ln(n-1) + \frac{2n-1}{n} \\ &= \frac{2(n^2-1)}{n} \left(\ln n + \ln\left(1 - \frac{1}{n}\right) \right) + \frac{2n-1}{n} \end{aligned}$$

We make use of the standard inequality stated below.

A Simple Example

Let the claim hold for all values up to $n - 1$. Then,

$$\begin{aligned} f(n) &= \frac{n+1}{n} f(n-1) + \frac{2n-1}{n} \\ &\leq \frac{n+1}{n} 2(n-1) \ln(n-1) + \frac{2n-1}{n} \text{ by induction hypothesis} \\ &= \frac{2(n^2-1)}{n} \ln(n-1) + \frac{2n-1}{n} \\ &= \frac{2(n^2-1)}{n} \left(\ln n + \ln\left(1 - \frac{1}{n}\right) \right) + \frac{2n-1}{n} \end{aligned}$$

We make use of the standard inequality stated below.

$$1 + x \leq e^x \text{ for } x \in R.$$

Hence,

$$\begin{aligned} f(n) &\leq \frac{2(n^2-1)}{n} \left(\ln n - \frac{1}{n} \right) + \frac{2n-1}{n} \\ &= 2n \ln n - \frac{2}{n} \ln n - 2 + \frac{2}{n^2} + 2 - \frac{1}{n} \\ &\leq 2n \ln n, \text{ establishing the inductive step.} \end{aligned}$$

A Simple Example

- Hence, the expected running time of the randomized quick sort algorithm is $O(n \ln n)$.
- But one of the limitations of the recurrence relation approach is that we do not know how the running time of the algorithm is spread around its expected value.
- Can this analysis be extended to answer questions such as, with what probability does the algorithm RandQuickSort need more than $12n \ln n$ time steps?
- Later on, we apply a different technique and establish that this probability is very small.
- To be able to answer such queries, we study Tail inequalities in the following.

Tail Inequalities

- We study three ways to estimate the tail probabilities of random variables.
- It will be noted that, **the more** information we know about the random variable **the better** the estimate we can derive about a given tail probability.

Tail Inequalities

- Markov Inequality: If X is a **non-negative valued** random variable with an expectation of μ , then $\Pr[X \geq c\mu] \leq 1/c$.
- Proof of Markov inequality:

$$\begin{aligned}\mu &= \sum_a a \cdot \Pr(X=a) \\ &= \sum_{a < c\mu} a \cdot \Pr(X=a) + \sum_{a \geq c\mu} a \cdot \Pr(X=a) \\ &\geq 0 + \sum_{a \geq c\mu} a \cdot \Pr(X=a) \\ &\geq c\mu \cdot \sum_{a \geq c\mu} \Pr(X=a) \\ &= c\mu \cdot \Pr(X \geq c\mu)\end{aligned}$$

Tail Inequalities

- Markov Inequality: If X is a non-negative valued random variable with an expectation of μ , then $\Pr[X \geq c\mu] \leq 1/c$.
- Applying this inequality tells us that the randomized quick sort algorithm has a run time of more than twice its expectation with a probability of $1/2$.
- The run time is n^2 with probability of nearly $\log n / n$.

Tail Inequalities

- Chebychev Inequality : We first define the terms standard deviation and variance of a random variable X .
- Let X be a random variable with an expectation of μ . The variance of X , denoted by $\text{var}(X)$, is defined as $\text{var}(X) = E[(X - \mu)^2]$. The standard deviation of X , denoted by σ_X , is defined as $\sigma_X = \sqrt{\text{var}(X)}$.
- Note that by definition, $\text{var}(X) = E[(X - \mu)^2] = E[X^2 - 2X\mu + \mu^2] = E[X^2] - \mu^2$.
- The second equality follows from the linearity of expectations.

Tail Inequalities

- Chebychev inequality: Let X be a random variable with expectation μ_X and standard deviation σ_X .
- Then, $\Pr[|X - \mu_X| \geq c\sigma_X] \leq 1/c^2$

Tail Inequalities

- Chebychev inequality: Let X be a random variable with expectation μ_X and standard deviation σ_X .
- Then, $\Pr[|X - \mu_X| \geq c\sigma_X] \leq \frac{1}{c^2}$
- Proof. Let random variable $Y = (X - \mu_X)^2$. Then,

$$E[Y] = E[(X - \mu_X)^2] = \sigma_X^2 \text{ by definition}$$

$$\begin{aligned} \text{Now, } \Pr[|X - \mu_X| \geq c\sigma_X] &= \Pr[(X - \mu_X)^2 \geq c^2\sigma_X^2] \\ &= \Pr[Y \geq c^2\sigma_X^2]. \end{aligned}$$

Applying Markov Inequality to the random variable Y

$$\Pr[Y \geq c^2\sigma_X^2] = \Pr[Y \geq c^2\mu_Y] \leq \frac{1}{c^2}$$

Tail Inequalities

- Better tail inequalities can be obtained by the powerful technique called Chernoff bounds.
- However, applicability is a little restricted too.
- Let us study the most popular version to start with.

Tail Inequalities

- Let X be a random variable defined as the sum of n **independent** and **identically distributed** random variables X_1, X_2, \dots, X_n .
 - In other words, $X = \sum_i X_i$
 - Short form i.i.d.
- Let us assume that each X_i is a Bernoulli random variable.
 - In other words, each X_i takes values in $\{0, 1\}$.
- Let $\Pr(X_i = 1) = p$ and hence $\Pr(X_i = 0) = 1 - p$.
- Finally, let $E[X] = \mu$.
 - Notice that $E[X] = \sum_i E[X_i] = n \cdot (1 \cdot p + (1-p) \cdot 0) = np$.

Tail Inequalities

- Several settings relate to the above statements.
- Consider throwing a (biased) coin over n trials.
- Each trial, the probability of Heads is p .
- So, each X_i corresponds to the fact that the i th trial results in a Heads.
- Let us count the number of Heads over the n trials. Indeed, $X = \sum_i X_i$ captures this count as a random variable.
- Note that the expected number of Heads over n trials is exactly np .

Tail Inequalities

- Finally, to the theorem.
- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- The normal strategy employed to prove tail estimates of sums of independent random variables is to make use of exponential moments.
- While proving Chebychev inequality, we made use of the second-order moment. It can be observed that using higher order moments would generally improve the bound on the tail inequality.
- But using exponential moments would result in a vast improvement.

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- While proving Chebychev inequality, we made use of the second-order moment of a random variable.
- It can be observed that using higher order moments would generally improve the bound on the tail inequality
- But using exponential moments would result in a vast improvement in the bound.
- An **exponential moment** of a random variable X is the expectation of functions of X such as e^X .

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- For each i , we define the random variable $Y_i = e^{tX_i}$ for a real number $t > 0$ that will be chosen later.
- Notice that
 - 1) Y_i is a positive valued random variable.

$$\begin{aligned} 2) E[Y_i] &= E[e^{tX_i}] = p_i \cdot e^t + (1-p_i) \cdot e^0 \\ &= p_i e^t + 1 - p_i \end{aligned}$$

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- Now define another random variable $Y = Y_1 \cdot Y_2 \cdots Y_n$.
- Now, we can note that

$$\begin{aligned} E[Y] &= E[Y_1 \cdot Y_2 \cdots Y_n] \\ &= \prod_{i=1}^n E[Y_i] = \left(p_i \cdot e^t + 1 - p_i \right)^n \end{aligned}$$

- Why does the above calculation hold?

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- The next step we do is do apply Markov inequality on Y as follows.
- First, notice that $Y = e^{tX}$ as

$$\begin{aligned} Y &= Y_1 \cdot Y_2 \cdot \dots \cdot Y_n = e^{tx_1} \cdot e^{tx_2} \cdot \dots \cdot e^{tx_n} \\ &= e^{t(x_1 + x_2 + \dots + x_n)} = e^{tX} \end{aligned}$$

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- First, notice that $Y = e^{tX}$. And,

- Further, $X \geq \mu(1+\delta) \Leftrightarrow e^{tX} \geq e^{t\mu(1+\delta)}$
 $\Rightarrow e^Y \geq e^{t\mu(1+\delta)}$

- So, we are interested in the event $Y \geq e^{t\mu(1+\delta)}$.

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- So, we are interested in the event $Y \geq t\mu(1+\delta)$. We proceed as:

$$\begin{aligned} \Pr(Y \geq e^{t\mu(1+\delta)}) &\leq \frac{E[Y]}{e^{t\mu(1+\delta)}} = \frac{\pi(1-p_i + p_i e^t)}{e^{t\mu(1+\delta)}} \\ &\leq \frac{\pi e^{-p_i + p_i e^t}}{e^{t\mu(1+\delta)}} \\ &= \frac{e^{-\mu(1-e^t)}}{e^{t\mu(1+\delta)}} \end{aligned}$$

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- So, we are interested in the event $Y \geq t\mu(1+\delta)$. We proceed as:

$$\begin{aligned} \Pr(Y \geq e^{t\mu(1+\delta)}) &\leq \frac{E[Y]}{e^{t\mu(1+\delta)}} = \frac{\pi(1-p_i + p_i e^t)}{e^{t\mu(1+\delta)}} \\ &\leq \frac{\pi e^{-p_i + p_i e^t}}{e^{t\mu(1+\delta)}} \\ &= \frac{e^{-\mu(1-e^t)}}{e^{t\mu(1+\delta)}} = e^{-\mu(1-e^t) - t\mu(1+\delta)} \end{aligned}$$

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- So, $\Pr(Y \geq e^{t\mu(1+\delta)}) = e^{-\mu(1-e^t) - t\mu(1+\delta)}$

- Since t is a free parameter in the above, we can find a t that minimizes the right hand side.

- To simplify, let $f(t) = \ln e^{-\mu(1-e^t) - t\mu(1+\delta)}$
 $= -\mu(1-e^t) - t\mu(1+\delta)$

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- To simplify, let $f(t) = \ln e^{-\mu(1-e^t) - t\mu(1+\delta)}$

$$= -\mu(1-e^t) - t\mu(1+\delta)$$

- Differentiating $f(t)$ wrt t , we get

$$f'(t) = \mu e^t - \mu(1+\delta)$$

- So, $f'(t) = 0$ at $t = \ln(1+\delta)$

- Verify that the above t corresponds to a minima.
OFFLINE.

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- With $t = \ln(1+\delta)$, we get that

$$\begin{aligned} \Pr(X \geq \mu(1+\delta)) &\leq \frac{e^{-\mu(1+\delta)t}}{(1+\delta)^{\mu(1+\delta)}} = \frac{e^{\mu\delta}}{(1+\delta)^{1+\delta}} \\ &= \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu \end{aligned}$$

- completing the proof.

Tail Inequalities

- Given the earlier conditions, it holds that for any $\delta > 0$,

$$\Pr(X \geq \mu(1+\delta)) \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- With $t = \ln(1+\delta)$, we get that

$$\Pr(X \geq \mu(1+\delta)) \leq \frac{e^{-\mu(1-(1+\delta))}}{(1+\delta)^{\mu(1+\delta)}} = \frac{e^{\mu\delta}}{(1+\delta)^{1+\delta}}$$

$$= \left(\frac{e^\delta}{(1+\delta)^{1+\delta}} \right)^\mu$$

- A simplification of the RHS gives

$$\Pr(X \geq \mu(1+\delta)) \leq \begin{cases} e^{-\mu\delta^2/4} & \text{if } \delta \leq 1 \\ e^{-\mu\delta \ln \delta} & \text{if } \delta > 1 \end{cases}$$

A Simple Example

- Here is one more practical application of the Chernoff bounds.
- Consider once again counting the number of heads out of tossing n fair coins independently.
 - So, $p = \frac{1}{2}$.
- Let X_i denote the random variable that takes 1 if the i th coin toss results in a head, and 0 otherwise.
 - $E[X_i] = \frac{1}{2}$.
- Let $X = \sum_i X_i$
- X counts the total number of heads over the n tosses.
 - $E[X] = n/2$.
 - With $n = 100$, say, we expect 50 heads over 100 coin tosses.

A Simple Example

- Markov inequality tells that the probability that X takes a value beyond 70 is $\Pr(X \geq 70) = \Pr(X \geq 70/50 \times 50) \leq 5/7 = 0.7$ (approx.)
- To apply Chebyshev's inequality, we need to do some extra work.
- $\text{Var}(X_i)$ for any i is computed as $E[X_i^2] - E[X_i]^2$.
- $E[X_i^2] = 0 \times (1/2) + 1 \times (1/2) = 1/2$.
- $\text{Var}(X_i) = 1/2 - (1/2)^2 = 1/4$.
- $\text{Var}(X) = 100 \times \text{Var}(X_1)$ due to independence.
- So, $\text{Var}(X) = 100/4 = 25$, and $\sigma_X = 5$.
- Now, $\Pr(X \geq 70)$ can be rewritten as $\Pr(|X - 50| \geq 20)$ and further as $\Pr(|X - 50| \geq (20 \times 1/\sigma_X) \sigma_X)$ which is now at most $25/400 = 1/16 = 0.0625$.

A Simple Example

- Markov inequality tells that the probability that X takes a value beyond 70 is $\Pr(X \geq 70) = \Pr(X \geq 70/50 \times 50) \leq 5/7$.
- Applying Chebychev's inequality, we get that $\Pr(X \geq 70)$ is at most $1/16$.
- This is quite an improvement on the upper bound of the probability of the event of interest.
- Let us see what Chernoff bounds will allow us to claim.
- We start by writing $\Pr(X \geq 70) = \Pr(X \geq (1+2/5) \times 50)$.
- From Chernoff bounds, this probability is at most $\exp\{-20\}/(1.4)^{70} \approx 0.028$ (approx.).

Applications of Tail Inequalities

- Here is one more practical application of the Chernoff bounds.
- Consider dividing a data set with features of interest into two parts: a test set and a training set.
- To make sure that both are roughly similar copies, you want to divide so that both data sets have the same number of data items of any given feature.
- Similar requirements arise in also experiments related to drug trials.

Example

Candidate	Above 50 Yrs	Diabetic	Frontline worker	Hypertension	Female
C1	1	0	1	1	0
C2	1	0	0	0	1
C3	0	0	0	0	1
C4	0	1	1	1	0
C5	1	1	0	0	0

Applications of Tail Inequalities

- To simplify matters, we will think of n data items with n features.
- Prepare a matrix A of $n \times n$ with entries from $\{0, 1\}$.
- Rows are features, columns are data items
- The goal is to find a vector x of size n with entries from $\{-1, +1\}$ such that Ax has the smallest possible maximum absolute entry.
 - Rows with $+1$ in x belong to one class and those with -1 in x belong to another class.
 - The maximum absolute entry in Ax indicates how many data items differ at feature i according to the division by x .

Applications of Tail Inequalities

- Example

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{bmatrix} \quad x = \begin{bmatrix} 1 \\ -1 \\ 1 \\ -1 \end{bmatrix} \quad Ax = \begin{bmatrix} -2 \\ 1 \\ 0 \\ -2 \end{bmatrix}$$

- The maximum absolute entry in Ax is 2.

Applications of Tail Inequalities

- No good deterministic algorithms are known.
- The brute-force algorithm has to try all possible 2^n vectors.
- However, a very simple randomized algorithm exists.
- Consider choosing each element of x uniformly at random from $\{1, -1\}$.
- We will show that the maximum absolute entry of Ax in such an x is bounded by $O((n \ln n)^{1/2})$ with high probability.

Applications of Tail Inequalities

- Consider choosing each element of x uniformly at random from $\{1, -1\}$.
- We will show that the maximum absolute entry of AX in such an X is bounded by $O((n \ln n)^{1/2})$ with high probability.
- Let the product $AX = Y$.
- Consider any Y_i , say Y_1 .
- By definition of matrix multiplication, $Y_1 = A_{11}X_1 + A_{12}X_2 + \dots + A_{1n}X_n$ where the A_{ij} denotes the element of A at the i th row and j th column.

Applications of Tail Inequalities

- Consider choosing each element of x uniformly at random from $\{1, -1\}$.
- We will show that the maximum absolute entry of AX in such an X is bounded by $O((n \ln n)^{1/2})$ with high probability.
- Let the product $AX = Y$.
- Consider any Y_i , say Y_1 .
- By definition of matrix multiplication, $Y_1 = A_{11}X_1 + A_{12}X_2 + \dots + A_{1n}X_n$
- Note that $E[X_i] = 0$ and by linearity of expectations, $E[Y_1] = 0$.

Applications of Tail Inequalities

- We will show that the maximum absolute entry of AX in such an X is bounded by $O((n \ln n)^{1/2})$ with high probability.
- Let the product $AX = Y$. Consider any Y_i , say Y_1 .
- By definition of matrix multiplication, $Y_1 = A_{11}X_1 + A_{12}X_2 + \dots + A_{1n}X_n$
- Note that $E[X_i] = 0$ and by linearity of expectations, $E[Y_1] = 0$.
- Let us imagine the case where the choices of X are made independently.
- We can then apply Chernoff bound on Y .

Applications of Tail Inequalities

- One small detour before we do that.
- In our Chernoff bound derived earlier, each X_i is $\{0,1\}$ valued random variable.
- Now, each X_i is a $+1/-1$ valued random variable.
- Need a new version of Chernoff bound!

Applications of Tail Inequalities

- Consider X as the random variable that is the sum of n independent and identically distributed random variables X_i with X_i taking value among $\{-1, +1\}$.
- Let $\Pr[X_i = +1] = \Pr[X_i = -1] = 1/2$.
- Note that $E[X] = n \cdot E[X_i] = n \cdot 0 = 0$.
- We want to know the $\text{Prob}(X \geq k)$ for some integral k .

Applications of Tail Inequalities

- Consider X as the random variable that is the sum of n independent and identically distributed random variables X_i with X_i taking value among $\{-1, +1\}$.
- Let $\Pr[X_i = +1] = \Pr[X_i = -1] = 1/2$.
- Note that $E[X] = n \cdot E[X_i] = n \cdot 0 = 0$.
- We want to know the $\text{Prob}(X \geq k)$ for some integral k .
- Instead of doing the entire calculation again, let us do the following.
- Define $Y_i = (1+X_i)/2$. Now, Y_i is $\{0,1\}$ valued.
- Define Y as the sum of Y_i 's.

Applications of Tail Inequalities

- Consider X as the random variable that is the sum of n independent and identically distributed random variables X_i with X_i taking value among $\{-1, +1\}$.
- We want to know the $\text{Prob}(X \geq k)$ for some integral k .
- Define $Y_i = (1+X_i)/2$. Now, Y_i is $\{0,1\}$ valued.
- Define Y as the sum of Y_i 's.
- Note that $EY = n/2$.
- Also, $X \geq k$ if and only if $Y \geq n/2 + k/2$.
- Now, $\Pr(X \geq k) = \Pr(Y \geq n/2 + k/2)$.

Applications of Tail Inequalities

- Also, $X \geq k$ if and only if $Y \geq n/2 + k/2$.
- Now, $\Pr(X \geq k) = \Pr(Y \geq n/2 + k/2)$.
- Rewrite as $\Pr(Y \geq E[Y](1+(k/n)))$ with $\delta = k/n < 1$.
- Applying Chernoff bounds with the above δ , we get that $\Pr(Y \geq E[Y](1+\delta)) \leq \exp(-E[Y]\delta^2/4) = \exp(-\delta^2/4n)$.

Applications of Tail Inequalities

- Back to our problem of set balancing....
- We now get $\Pr(Y_1 \geq 8 \sqrt{n \ln n}) \leq \exp(-64n \ln n) = \exp(-8 \ln n) = 1/n^8$.
- But we are interested in a two-sided bound.
- That is, since we want to minimize the absolute value of Y_1 , we need to compute $\Pr[Y_1 \leq -d]$ also.
- But by symmetry, we have that $\Pr(Y_1 \leq -d) \leq 1/n^8$.
- So, $\Pr[|Y_1| \geq (8n \ln n)^{1/2}] \leq 2/n^8$.

Applications of Tail Inequalities

- Back to our problem of set balancing....
- So, $\Pr[|Y_1| \geq (8n \ln n)^{1/2}] \leq 2/n^8$.
- But, what about Y_2, Y_3 , etc.
- This is where another classical probability result aids us.
- **Boole's inequality**. For any events, E_1, E_2, \dots, E_n ,
- $\Pr(E_1 \cup E_2 \cup \dots \cup E_n) \leq \Pr(E_1) + \Pr(E_2) + \dots + \Pr(E_n)$.
- Apply the above to get that with probability at least $1 - 2/n^7$, every Y_i has an absolute value that is within $(8n \ln n)^{1/2}$.

Applications of Randomization

- It is known that randomization combined with (non-trivial) algebraic techniques can lead to important applications.
- In this section, we showcase some of such techniques with respect to verification of identities.
- One such technique is called the fingerprinting technique described as follows.

Applications of Randomization

- Let U be any universe of objects and x and y be any two elements from U .
- The question we ask is, Is $x = y$?
- One can answer this using at least $\log |U|$ bits in a deterministic manner.
- However, consider mapping elements of U to a sparse universe V such that x and y are identical if and only if their images in V are identical, with a good chance.
- These images can be thought of as fingerprints of x and y .

Applications of Randomization

- Let us apply the above technique to matrix product verification. Let F be a field and A and B are two matrices with entries from F .
- Suppose it is claimed that $C = A \cdot B$.
- The fastest known matrix multiplication algorithm runs in time $O(n^{2.376})$.
- This algorithm is very difficult to implement, but the standard algorithms such as the Strassen's recursive algorithm takes time $O(n^{\log_2 7})$.
- So, to verify if C is indeed the product of A and B , it takes time equal to multiplying two matrices.

Applications of Randomization

- However, a simpler and efficient randomized approach exists.
- Let r be any vector with entries being 0 or 1.
- Let each element of r be chosen independently and uniformly at random.
- It is being assumed without loss of generality that 0 and 1 are the additive and multiplicative identities of the field F .

Applications of Randomization

- Compute $x = Br$, and $y = Ax$.
- Similarly, compute $z = Cr$.
- If $A \cdot B = C$ is indeed true, then y must equal z for any r .
- Also, x , y , and z can each be computed in $O(n^2)$ time.
- So, the time efficiency is established.
- It remains to see the verification efficiency.
- The following lemma argues that the verification procedure is efficient.

Applications of Randomization

- Lemma:1 Let A , B , and C be $n \times n$ matrices from F such that $AB \neq C$. Then, for r chosen uniformly at random from $\{0, 1\}^n$, $\Pr(ABr = Cr) \leq 1/2$.
- Proof. Consider the matrix $D := AB - C$. Since, $AB \neq C$, matrix D is not the matrix of all zeros.
- We are interested in the event that $Dr = 0$.
- Assume without loss of generality that the first row of D has a nonzero entry and all nonzero entries in that row are before any zero entry.

Applications of Randomization

- Lemma:1 Let A , B , and C be $n \times n$ matrices from F such that $A \neq B$. Then, for r chosen uniformly at random from $\{0, 1\}^n$, $\Pr(Ar = Br) \leq 1/2$.
- Proof. Consider the first row of A and the scalar obtained by multiplying the first row of A with r .
- The result is zero if and only if:
$$r_1 = -\frac{\sum_{i=1}^k D_{1i} r_i}{D_{11}}$$
- In the above, it is assumed that there are $k > 0$ nonzero elements in the first row of A .

Applications of Randomization

- Lemma:1 Let A , B , and C be $n \times n$ matrices from F such that $Ax \neq Cx$. Then, for r chosen uniformly at random from $\{0, 1\}^n$, $\Pr(Ar = Cr) \leq 1/2$.
- Proof. The above event is a super-event of the event that $Dr = 0$.
- Therefore the probability of the event $Dr = 0$ is upper bounded by the probability of the above event.
- To compute the above probability, imagine that all the choices r_2, \dots, r_k have been made.
- In that case, the right hand side is a scalar from the field F .
- The left hand side is a value uniformly chosen amongst (at least) two values in F . The required probability therefore cannot exceed $1/2$.

Applications of Randomization

- To improve the verification efficiency of the procedure, we can also use repeated independent trials.
- Let us perform t independent trials of the above procedure.
- For $AB \neq C$, the probability that the test fails in each trial is at most $1/2$.
- So, in t trials, the probability that all t trials fail, or a majority of the t trials fail, is at most $(1/2)^t$.
- For $t = O(\log n)$, the failure probability is polynomially small

Applications of Randomization

- Other applications of the technique include verifying polynomial identities.
- For instance, let $P_1(x)$, $P_2(x)$ be two polynomials in a field F .
- The polynomial product verification problem is to check whether $P_1(x) \cdot P_2(x) = P_3(x)$ for a given $P_3(x)$.
- It holds that there exists an $O(n \log n)$ time algorithm to multiply two polynomials, where n is the maximum degree of P_1 and P_2 .
- We design a verification algorithm that is faster than $O(n \log n)$.

Applications of Randomization

- Let $S \subset F$ be a subset of size at least $2n + 1$.
- The main idea of the verification procedure is that if indeed $P_3(x)$ equals $P_1(x) \cdot P_2(x)$, then, also $P_3(r) = P_1(r) \cdot P_2(r)$ for r chosen uniformly at random from S .
- Further, evaluating a polynomial at a given input can be done in $O(n)$ time.
- So, we can declare that $P_3(x)$ equals $P_1(x) \cdot P_2(x)$ unless $P_3(r) \neq P_1(r) \cdot P_2(r)$.
- The algorithm makes a mistake only when indeed $P_3(x) \neq P_1(x) \cdot P_2(x)$ but the choice of r fails to detect this.

Applications of Randomization

- To estimate the probability that the algorithm makes a mistake, let $Q(x) := P_3(x) - P_1(x) \cdot P_2(x)$.
- The degree of $Q(x)$ is at most $2n$.
- Suppose that $P_3(x) \neq P_1(x) \cdot P_2(x)$.
- Then, $Q(x)$ is a nonzero polynomial.
- So the test fails if $Q(r) = 0$.
- However, the polynomial $Q(x)$ of degree at most $2n$ can have at most $2n$ roots.
- So, the probability that $Q(r) = 0$ is at most $2n/|S|$, which is also the probability of error.
- As earlier, the probability of failure can be made polynomially small in n by using repeated trials or choosing a larger S , or both.

Applications of Randomization

- One may wonder whether it is at all worthwhile to have elaborate verification algorithms for things as simple as polynomial product verification.
- Such techniques however are more applicable when polynomials are not available explicitly.

Finger-Printing

- The two algorithms that we considered today have the property that for inputs that are identical, the algorithm does not make any error.
- But in inputs that are not identical, the algorithm makes an error that is upper bounded by at least a constant.
- Repeated runs of the same algorithm can catch the error, hence the error can be made arbitrarily small.

Finger-Printing

- On inputs that are not identical, the algorithm makes an error that is upper bounded by at least a constant.
- Repeated runs of the same algorithm can catch the error, hence the error can be made arbitrarily small.
- Consider A to be a finger-printing algorithm.
- Let us run A on input x, y for t iterations.
- The t outputs are, say, o_1, o_2, \dots, o_t .
- If any of these t outputs are NO, then we can return NO as the answer.
- If all t are YES, then we return YES as the answer.
- Given that $x = y$,

Finger-Printing

- Repeated runs of the same algorithm can catch the error, hence the error can be made arbitrarily small.
- Consider A to be a finger-printing algorithm.
- Let us run A on input x, y for t iterations.
- The t outputs are, say, o_1, o_2, \dots, o_t .
- If any of these t outputs are NO, then we can return NO as the answer.
- If all t are YES, then we return YES as the answer.
- Given that $x = y$, $\Pr(A(x,y) = \text{YES}) = 1$.
- Given that $x \neq y$, $\Pr(A(x,y) = \text{YES}) \leq (1/2)^t$.
- So, if $t = O(\log n)$, then the error probability is $O(1/n^c)$.

RP and co-RP

- The above algorithms are called as co-RP algorithms.
- RP stands for Randomized Polynomial.
- Definition: The class RP consists of languages L such that there exists a randomized algorithm running in worst case polynomial time such that for any input x :
 - $x \in L \Rightarrow \Pr(A \text{ accepts } x) \geq 1/2$.
 - $x \notin L \Rightarrow \Pr(A \text{ accepts } x) = 0$
- The complement of the class RP is the class co-RP.
- Note that any RP or co-RP algorithm can err only on one side, either for x in L , or for x not in L .

RP and co-RP

- In definition: The class RP consists of languages L such that there exists a randomized algorithm running in worst case polynomial time such that for any input x :
 - $x \in L \Rightarrow \Pr(A \text{ accepts } x) \geq 1/2$.
 - $x \notin L \Rightarrow \Pr(A \text{ accepts } x) = 0$
- The complement of the class RP is the class co-RP.
- In definition: The class co-RP consists of languages L such that there exists a randomized algorithm running in worst case polynomial time such that for any input x :
 - $x \in L \Rightarrow \Pr(A \text{ accepts } x) = 0$.
 - $x \notin L \Rightarrow \Pr(A \text{ accepts } x) \leq 1/2$.

RP and co-RP

- In definition: The class RP consists of languages L such that there exists a randomized algorithm running in worst case polynomial time such that for any input x :
 - $x \in L \Rightarrow \Pr(A \text{ accepts } x) \geq 1/2$.
 - $x \notin L \Rightarrow \Pr(A \text{ accepts } x) = 0$
- In definition: The class co-RP consists of languages L such that there exists a randomized algorithm running in worst case polynomial time such that for any input x :
 - $x \in L \Rightarrow \Pr(A \text{ accepts } x) = 0$.
 - $x \notin L \Rightarrow \Pr(A \text{ accepts } x) \leq 1/2$
- Note that any RP or co-RP algorithm can err only on one side, either for x in L , or for x not in L .

RP and co-RP

- In definition: The class RP consists of languages L such that there exists a randomized algorithm running in worst case polynomial time such that for any input x :
 - $x \in L \Rightarrow \Pr(A \text{ accepts } x) \geq 1/2$.
 - $x \notin L \Rightarrow \Pr(A \text{ accepts } x) = 0$
- Can you recall any RP algorithms that we studied?
-

RP and co-RP

- In definition: The class RP consists of languages L such that there exists a randomized algorithm running in worst case polynomial time such that for any input x :
 - $x \in L \Rightarrow \Pr(A \text{ accepts } x) \geq 1/2$.
 - $x \notin L \Rightarrow \Pr(A \text{ accepts } x) = 0$
- Can you recall any RP algorithms that we studied?
- The interactive proof algorithm for graph non-isomorphism is an example.

What about Randomized QuickSort

- The randomized quick sort algorithm does not make any errors in its output.
- So, clearly, it does not fit into either of RP or co-RP.
- While there are no errors in its output, we recall that its run time may vary.
- Such algorithms are called ZP algorithms.
 - Stands for Zero Error Expected Polynomial, time
- The class ZP consists of languages L such that there is a randomized algorithm A that **always** outputs the **correct** answer while running in expected polynomial time.
- Another name for ZP algorithms is Las Vegas algorithms.

Proof by Existence

- Many times you want to show that a particular combinatorial object exists.
- May be very inefficient to build possibly because of a huge space and a small target of interest.
 - Like finding a needle in a haystack.
- This is where randomization can come to the rescue.

Proof by Existence

- Two useful statements:
 - 1) If a random variable has expected value $E[X] = a$, then certainly there exists a realisation of X with value $\geq a$ and a realisation of X with value $\leq a$.
 - 2) If a random object drawn from some universe of objects has a certain property with non-zero probability then there must exist an object with that property in this universe.

Proof by Existence

- Two useful statements:
 - 1) If a random variable has expected value $E[X] = a$, then certainly there exists a realisation of X with value $\geq a$ and a realisation of X with value $\leq a$.
 - 2) If a random object drawn from some universe of objects has a certain property with non-zero probability then there must exist an object with that property in this universe.
- Each of these statements looks simple on their own, but they are remarkably powerful in Computer Science.

Proof by Existence

- We will start with a simple example.
- Consider an undirected graph $G = (V, E)$.
- We want to find a subgraph G' of G that:
 - 1) has the largest number of edges of G , and
 - 2) G' is bipartite.
- We will first show the existence of a G' with $|E(G')|$ at least $|E(G)|/2$.

Proof by Existence

- Consider an undirected graph $G = (V, E)$.
- We want to find a subgraph G' of G that:
 - 1) has the largest number of edges of G , and
 - 2) G' is bipartite.
- We will first show the existence of a G' with $|E(G')|$ at least $|E(G)|/2$.
- The random experiment we perform is to assign a bit 0 or 1, denoted $b(v)$, to each vertex v of G ind. and uar.
- Put all vertices in G' .
- An edge uv is in G' if and only if $b(u)$ is different from $b(v)$.

Proof by Existence

- The random experiment we perform is to assign a bit 0 or 1, denoted $b(v)$, to each vertex v of G ind. and uar.
- Put all vertices in G' .
- An edge uv is in G' if and only if $b(u)$ is different from $b(v)$.
- In other words, $G' = (V_0 \cup V_1, V_0 \times V_1 \cap E(G))$.
- Let us now bound $|E(G')|$.
- Let X_{uv} be a random variable that indicates the event $\{uv \in E(G')\}$.
- $E[X_{uv}] = \Pr(\{uv \in E(G')\}) = 1/2$.

Proof by Existence

- The random experiment we perform is to assign a bit 0 or 1, denoted $b(v)$, to each vertex v of G ind. and uar.
- An edge uv is in G' if and only if $b(u)$ is different from $b(v)$.
- Let us now bound $|E(G')|$.
- Let X_{uv} be a random variable that indicates the event $\{uv \text{ in } E(G')\}$.
- $E[X_{uv}] = \Pr(\{uv \text{ in } E(G')\}) = 1/2$.
- Let $X = \sum_{uv \text{ in } E(G)} X_{uv}$.
- $E[X] = E[\sum_{uv \text{ in } E(G)} X_{uv}] = \sum_{uv \text{ in } E(G)} E[X_{uv}] = |E(G)|/2$.

Proof by Existence

- The random experiment we perform is to assign a bit 0 or 1, denoted $b(v)$, to each vertex v of G ind. and uar.
- Let X_{uv} be a random variable that indicates the event $\{uv \in E(G')\}$.
- $E[X_{uv}] = \Pr(\{uv \in E(G')\}) = 1/2$.
- Let $X = \sum_{uv \in E(G)} X_{uv}$.
- $E[X] = E[\sum_{uv \in E(G)} X_{uv}] = \sum_{uv \in E(G)} E[X_{uv}] = |E(G)|/2$.
- By Statement (1) earlier, there must exist an assignment of $b()$ to vertices such that G' has at least half the edges of G .

Proof by Existence

- Let us now move to a more involved example.
- We start with the definition of an (α, β, n, d) expander.
- A bipartite graph $G = (V_1 \cup V_2, E)$ on n nodes is an (α, β, n, d) expander if
 - 1) Every vertex in V_1 has degree at most d .
 - 2) For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- Ideally, d should be small and β as large as possible.

Proof by Existence

- Let us now move to a more involved example.
- We start with the definition of an (α, β, n, d) expander.
- A bipartite graph $G = (V_1 \cup V_2, E)$ on n nodes is an (α, β, n, d) expander if
 - 1) Every vertex in V_1 has degree at most d .
 - 2) For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- To build such a graph in a deterministic manner is not easy.
- Simple randomized construction with $d = 18$, $\alpha = 1/3$, and $\beta = 2$

Proof by Existence

- A bipartite graph $G = (V_1 \cup V_2, E)$ on n nodes is an (α, β, n, d) expander if
 - 1) Every vertex in V_1 has degree at most d .
 - 2) For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- Simple randomized construction with $d = 18$, $\alpha = 1/3$, and $\beta = 2$.
- We will actually not use these values until the very end of the proof.

Proof by Existence

- A bipartite graph $G = (V_1 \cup V_2, E)$ on n nodes is an (α, β, n, d) expander if
 - 1) Every vertex in V_1 has degree at most d .
 - 2) For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- Let each vertex v in V_1 choose d neighbors in V_2 by sampling independently and uniformly at random.
- We can even sample with replacement.
- In other words, the same choice can be made more than once.
- We will still consider only one copy of any multiple choices.

Proof by Existence

- A bipartite graph $G = (V_1 \cup V_2, E)$ on n nodes is an (α, β, n, d) expander if
 - 1) Every vertex in V_1 has degree at most d .
- Let each vertex v in V_1 choose d neighbors in V_2 by sampling independently and uniformly at random.
- By this construction, each vertex in V_1 has degree at most d .
- Next, we show the second condition.

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- Let $|V_1| = |V_2| = n$.
- Let each vertex v in V_1 choose d neighbors in V_2 by sampling independently and uniformly at random.
- Consider any subset S of V_1 with $|S| = s$.
- Let T be any subset of V_2 of size $< \beta s$.
- Consider the event that all the neighbors of vertices in S are in T .
- This event occurs with probability $(\beta s/n)^{ds}$.

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- Let $|V_1| = |V_2| = n$. Let each vertex v in V_1 choose d neighbors in V_2 by sampling independently and uar.
- Consider any subset S of V_1 with $|S| = s$.
- Let T be any subset of V_2 of size $< \beta s$.
- Consider the event that all the neighbors of vertices in S are in T .
- This event occurs with probability at most $(\beta s/n)^{ds}$.
- We have to now look at all possible S and all possible T .

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- Consider the event that all the d_S neighbors of vertices in S of size s are in T .
- This event occurs with probability at most $(\beta s/n)^{d_S}$.
- Let us use Boole's inequality to upper bound the probability of the event that for some S all its neighbors are in T .
- We have to now look at all possible S and all possible T .
- There are $\binom{n}{s}$ ways to choose S and $\binom{n}{\beta s}$ ways to choose T .

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- This event occurs with probability at most $(\beta s/n)^{ds}$.
- There are $\binom{n}{s}$ ways to choose S and $\binom{n}{\beta s}$ ways to choose T .
- The probability that for some S all its neighbors are in T is now upper bounded by $\binom{n}{s} \cdot \binom{n}{\beta s} \cdot (\beta s/n)^{ds}$.
- To simplify, let us use the inequality that for any n and k , $\binom{n}{k}$ is at most $(en/k)^k$.

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- The probability that for some S all its neighbors are in T is now upper bounded by $\binom{n}{s} \cdot \binom{n}{\beta s} \cdot (\beta s/n)^{ds}$.
- To simplify, let us use the inequality that for any n and k , $\binom{n}{k}$ is at most $(en/k)^k$.
- The probability is at most $(en/s)^s \cdot (en/\beta s)^{\beta s} \cdot (\beta s/n)^{ds}$.
- Simplifying we get, $\left[(s/n)^{d-\beta-1} e^{1+\beta} \beta^{d-\beta} \right]^s$.
- Use that s is at most αn , for $\alpha = 1/3$ to simplify to:

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- The probability that for some S all its neighbors are in T is now upper bounded by $\binom{n}{s} \cdot \binom{n}{\beta s} \cdot (\beta s/n)^{ds}$.
- The probability is at most $(en/s)^s \cdot (en/\beta s)^{\beta s} \cdot (\beta s/n)^{ds}$.
- Simplifying we get, $\left[(s/n)^{d-\beta-1} e^{1+\beta} \beta^{d-\beta} \right]^s$.
- Use that s is at most αn , for $\alpha = 1/3$ to simplify the above to $\left[(\beta/3)^d (3e)^{1+\beta} \right]^s$.
- Use $d = 18$ and $\beta = 2$ to simplify to $\left[(2/3)^{18} (3e)^3 \right]^s$.

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- The probability that for some S all its neighbors are in T is at most $(en/s)^s \cdot (en/\beta s)^{\beta s} \cdot (\beta s/n)^{ds}$.
- Use that s is at most αn , for $\alpha = 1/3$ to simplify the above to $\left[(\beta/3)^d (3e)^{1+\beta} \right]^s$.
- Use $d = 18$ and $\beta = 2$ to simplify to $\left[(2/3)^{18} (3e)^3 \right]^s$.
- Notice that the term in $[]$ is at most $1/2$. So, the entire probability is at most $(1/2)^s$.

Proof by Existence

- Condition 2: For any subset S of vertices from V_1 such that $|S|$ is at most αn , there are at least $\beta|S|$ neighbors in V_2 .
- The probability that for some S all its neighbors are in T is at most $\left[(2/3)^{18} (3e)^3 \right]^s$.
- Notice that the term in $[]$ is at most $1/2$. So, the entire probability is at most $(1/2)^s$.
- We used a specific s . But, we need to show the result for all s between 1 to αn .
- Apply Boole's inequality again to get that

$$\begin{aligned} & \sum_{s > 0} \Pr(\text{for some } S \text{ all its neighbors are in } T) \\ & \leq \sum_{s > 0} (1/2)^s < 1 \end{aligned}$$

Proof by Existence--Another Example

- Consider the following claim.
- There is a bipartite graph $G = (L, R, E)$ such that
 - $|L| = n$
 - $|R| = 2^{\log^2 n}$
 - Every subset of $n/2$ vertices of L has at least $2^{\log^2 n} - n$ neighbors in R .
 - No vertex of R has more than $12\log^2 n$ neighbors.
- We want to use the technique of proof by existence to show the above claim.

Proof by Existence--Another Example

- There is a bipartite graph $G = (L, R, E)$ such that
 - $|L| = n$, $|R| = 2^{\log^2 n}$. Every subset of $n/2$ vertices of L has at least $2^{\log^2 n} - n$ neighbors in R . No vertex of R has more than $12\log^2 n$ neighbors.
 - Let every vertex of L choose d neighbors in R independently and uniformly at random.
 - Choices are made with replacement.
 - Multiple edges are dropped in favor of one edge.

Proof by Existence--Another Example

- There is a bipartite graph $G = (L, R, E)$ such that
 - $|L| = n$, $|R| = 2^{\log^2 n}$. Every subset of $n/2$ vertices of L has at least $2^{\log^2 n} - n$ neighbors in R . No vertex of R has more than $12\log^2 n$ neighbors.
 - Let every vertex of L choose d neighbors in R independently and uniformly at random.
 - Let us now estimate the degree of any vertex of R .
 - Let $|R| = r$.
 - We can think of the degree of a vertex v in R as the expectation of the random variable X that indicates how many vertices in L choose v as a neighbor.
 - Each vertex in L makes d choices, so we have nd choices in all.

Proof by Existence--Another Example

- There is a bipartite graph $G = (L, R, E)$ such that
 - $|L| = n$, $|R| = 2^{\log^2 n}$. Every subset of $n/2$ vertices of L has at least $2^{\log^2 n} - n$ neighbors in R . No vertex of R has more than $12\log^2 n$ neighbors.
 - Let every vertex of L choose d neighbors in R independently and uniformly at random.
 - Let $|R| = r$.
 - We can think of the degree of a vertex v in R as the expectation of the random variable X that indicates how many vertices in L choose v as a neighbor.
 - Each neighbor in L makes d choices, so we have nd choices in all.
 - Let X_i be a random variable if the i th choice is v .

Proof by Existence--Another Example

- Let $|R| = r$.
- We can think of the degree of a vertex v in R as the expectation of the random variable X that indicates how many vertices in L choose v as a neighbor.
- Each neighbor in L makes d choices, so we have nd choices in all.
- Let X_i be a random variable if the i th choice is v .
- $E[X_i] = 1/r$.
- $X = \sum X_i$ and so $E[X] = \sum E[X_i] = nd/r$.
- Pick $d = r \cdot 2 \log^2 n / n$ so that $E[X] = 2 \log^2 n$.
- Now apply Chernoff bounds on X for the event $X \geq 12 \log^2 n$.
- Use Boole's inequality to bound the probability of the bad event for every v in R .

Proof by Existence--Another Example

- There is a bipartite graph $G = (L, R, E)$ such that
 - $|L| = n$, $|R| = 2^{\log^2 n}$. Every subset of $n/2$ vertices of L has at least $2^{\log^2 n} - n$ neighbors in R .
 - Let every vertex of L choose d neighbors in R independently and uniformly at random.
 - We now move to property 1.
 - Let S be any subset of size $n/2$ from L .
 - Let T be any subset of R of size $2^{\log^2 n} - n$.
 - Consider the event that all the neighbors of S are in T .
 - This happens with a probability of $\left[(2^{\log^2 n} - n)/n\right]^{nd/2}$.

Proof by Existence--Another Example

- There is a bipartite graph $G = (L, R, E)$ such that
 - $|L| = n$, $|R| = 2^{\log^2 n}$. Every subset of $n/2$ vertices of L has at least $2^{\log^2 n} - n$ neighbors in R .
 - Let S be any subset of size $n/2$ from L .
 - Let T be any subset of R of size $2^{\log^2 n} - n$.
 - Consider the event that all the neighbors of S are in T .
 - This happens with a probability of $[(r - n)/r]^{nd/2}$.
 - Now, consider all possible choices of S and T . The probability that for any S all its neighbors are in some T is upper bounded by: $\binom{n}{n/2} \cdot \binom{r}{r-n} \cdot [(r - n)/r]^{nd/2}$.
 - We will now show that the above probability is at most 1.

Proof by Existence--Another Example

- Now, consider all possible choices of S and T . The probability that for any S all its neighbors are in some T is upper bounded by: $\binom{n}{n/2} \cdot \binom{r}{r-n} \cdot [(r-n)/r]^{nd/2}$.
- We will now show that the above probability is at most 1.
- Use the (in)equalities
 - $\binom{n}{n-k} = \binom{n}{k}$ for k between 0 and n .
 - $\binom{n}{k}$ is at most $(en/k)^k$.
 - $(1+x)$ is at most e^x for any real number x .
- The required probability is
 - $(2e)^{n/2} \cdot (er/n)^n \cdot (e)^{-n^2d/2r}$.
 - Recall that $d = 2\log^2 n \cdot r/n$.

Proof by Existence--Another Example

- Now, consider all possible choices of S and T . The probability that for any S all its neighbors are in some T is upper bounded by: $\binom{n}{n/2} \cdot \binom{r}{r-n} \cdot [(r-n)/r]^{nd/2}$.
- We will now show that the above probability is at most 1.
- Use the (in)equalities
 - $\binom{n}{n-k} = \binom{n}{k}$ for k between 0 and n .
 - $\binom{n}{k}$ is at most $(en/k)^k$.
 - $(1+x)$ is at most e^x for any real number x .
- The required probability is
 - $(2e)^{n/2} \cdot (er/n)^n \cdot (e)^{-n^2 d/2r}$.
- Recall that $d = 2\log^2 n \cdot r/n$ and $\log r = \log^2 n$ to simplify.

Yet Another Example

- Consider a Boolean formula in CNF.
- In CNF, each clause is a disjunction of literals
- The formula is a conjunction of clauses.
- Another name for CNF is Product-of-Sums.

Yet Another Example

- We show that for any set of m clauses, there is a truth assignment that satisfies at least $m/2$ clauses.
- Proof: Consider a random assignment of truth values to variables as T/F.
- Consider a clause C_i of k variables.
- C_i is not satisfied with probability 2^{-k} .
- Define a random variable Z_i that indicates the event C_i is satisfied.
- $E[Z_i] = \Pr(C_i \text{ is satisfied}) = 1 - 2^{-k}$.
- Define Z as the number of clauses satisfied. $Z = \sum Z_i$.
- $E[Z] = E[\sum Z_i] = \sum E[Z_i] = m(1 - 2^{-k}) \geq m/2$ as $k \geq 1$.

Yet Another Example

- We show that for any set of m clauses, there is a truth assignment that satisfies at least $m/2$ clauses.
- Proof: Consider a random assignment of truth values to variables as T/F.
- $E[Z] = E[\sum Z_i] = \sum E[Z_i] = m(1 - 2^{-k}) \geq m/2$ as $k \geq 1$.
- The above holds irrespective of whether the formula is satisfiable or not.
- This version of the problem where we intend to maximize the number of clauses that can be satisfied is called as MAXSAT.
- MAXSAT is also NP-hard indicating that no good polynomial solutions exist.

Yet Another Example

- This version of the problem where we intend to maximize the number of clauses that can be satisfied is called as MAXSAT.
- Define for an instance I , $m^*(I)$ to be the maximum number of clauses that can be satisfied.
- Let $m^A(I)$ be the number of (expected) clauses that can be satisfied by an (randomized) algorithm A .
- The ratio $m^A(I)/m^*(I)$ is the performance ratio of algorithm A .
- We seek algorithms that this ratio as close to 1.
- The previous approach gives us $1/2$ as the ratio.

Yet Another Example

- This version of the problem where we intend to maximize the number of clauses that can be satisfied is called as MAXSAT.
- The ratio $m^A(I)/m^*(I)$ is the performance ratio of algorithm A.
- We seek algorithms that this ratio as close to 1.
- The previous approach gives us $1/2$ as the ratio.
 - Actually the ratio is $1-2^{-k}$.
- In fact, there are instances where one can satisfy only $1/2$ of the clauses.

LP Rounding

- This version of the problem where we intend to maximize the number of clauses that can be satisfied is called as MAXSAT.
- We now study an approach that does better than $1/2$.
- Finally, we devise an algorithm that gets us a ratio of $3/4$.

LP Rounding

- The technique of LP Rounding uses the following approach.
- Write the optimization problem as an integer linear program (ILP).
- Relax some of the constraints of the ILP in a step called LP Relaxation to convert the ILP to a simple Linear Program (LP).
- Note that LP can be solved in polynomial time. Get an optimal solution to the LP.
- Round the solution from LP to satisfy the integrality constraints.
- May lose some quality in this step but that is inevitable.

LP Rounding

- Let us apply LP rounding to the MAXSAT problem.
- Consider a clause C_i .
- An indicator variable z_i with values in $\{0, 1\}$ is defined to indicate whether C_i is satisfied or not.
- We now seek to maximize $\sum_i z_i$.
- For each variable x_j , we define an indicator variable y_j that takes values 1 or 0 corresponding to $x_j = \text{True}$ or False respectively.
- Since variables can appear in either the pure form or the complemented form, we separate these as follows.

LP Rounding

- Let us apply LP rounding to the MAXSAT problem.
- Consider a clause C_i .
- For each variable x_j , we define an indicator variable y_j that takes values 1 or 0 corresponding to $x_j = \text{True}$ or False respectively.
- Since variables can appear in either the pure form or the complemented form, we separate these as follows.
- Define C_{i+} to be the indices of variables that appear in pure form in C_i .
- Define C_{i-} to be the indices of variables that appear in pure form in C_i .

LP Rounding

- Let us apply LP rounding to the MAXSAT problem.
- Consider a clause C_i .
- For each variable x_j , we define an indicator variable y_j that takes values 1 or 0 corresponding to $x_j = \text{True}$ or False respectively.
- Define C_{i+} to be the indices of variables that appear in pure form in C_i .
- Define C_{i-} to be the indices of variables that appear in pure form in C_i .
- Now, clause C_i is satisfied if it holds that for each i

$$\sum_{j \in C_{i+}} y_j + \sum_{j \in C_{i-}} (1 - y_j) \geq z_i.$$

LP Rounding

- Let us apply LP rounding to the MAXSAT problem.
- The entire integer linear program is:

Maximize $\sum_i z_i$

subject to

$$\sum_{j \in C_{i+}} y_j + \sum_{j \in C_{i-}} (1 - y_j) \geq z_i \text{ for all } i.$$

where $y_j, z_i \in \{0, 1\}$ for all i and j .

LP Rounding

- Let us apply LP rounding to the MAXSAT problem.
- Let us relax the constraints on y_j and z_i so that they can take values in $[0,1]$
- Note they are real numbers between 0 and 1 now and not just integral necessarily.
- We will use u_j and v_i for the values of the best solution to the relaxed linear program.
 - We use u 's for the variables and v 's for the clauses.
- Notice that $\sum_i v_i$ is an upper bound on the number of clauses that can be satisfied.
- But, the values of u_j are not integral, so they do not yet correspond to True/False values in a truth assignment.

LP Rounding

- Let us relax the constraints on y_j and z_i so that they can take values in $[0,1]$
- We will use u_j and v_i for the values of the best solution to the relaxed linear program.
- But, the values of u_j are not integral, so they do not yet correspond to True/False values in a truth assignment.
- The next step in the technique suggests to round the u_j 's so that a truth assignment can be obtained. This step is called randomized rounding.
- Our rounding does the following: Set y_j to 1 with probability u_j .
 - This sets x_j to True with the same probability.

LP Rounding

- The next step in the technique suggests to round the u_j 's so that a truth assignment can be obtained. This step is called randomized rounding.
- Our rounding does the following: Set y_j to 1 with probability u_j .
 - This sets x_j to True with the same probability.
- We now estimate the probability that a clause C_i is satisfied.

LP Rounding

- We now estimate the probability that a clause C_i is satisfied.
- Claim: A clause C_i with k literals is satisfied with probability at least $1 - (1 - 1/k)^k \cdot v_i$.
- Recall what is v_i .
- Let us assume wlog that all the variables in C_i appear in their pure form.
- So, $C_i = x_1 \vee x_2 \vee \dots \vee x_k$ for some variables x_1 through x_k .
- In the relaxed LP, we satisfy the constraint $u_1 + u_2 + \dots + u_k \geq v_i$.
- C_i now remains unsatisfied if the corresponding x_1 through x_k are all 0.

LP Rounding

- A clause C_i with k literals is satisfied with probability at least $1 - (1-1/k)^k \cdot v_i$.
- Recall what is v_i .
- So, $C_i = x_1 \vee x_2 \vee \dots \vee x_k$ for some variables x_1 through x_k .
- In the relaxed LP, we satisfy the constraint $u_1 + u_2 + \dots + u_k \geq v_i$.
- C_i now remains unsatisfied if the corresponding x_1 through x_k are all 0.
- This happens with probability $(1-u_j)$ for each variable and hence with probability $\prod_j (1-u_j)$ for the k variables.
- So, C_i is satisfied with probability $1 - \prod_j (1-u_j)$.

LP Rounding

- A clause C_i with k literals is satisfied with probability at least $1 - (1-1/k)^k \cdot v_i$.
- This happens with probability $(1-u_j)$ for each variable and hence with probability $\prod_j (1-u_j)$ for the k variables.
- So, C_i is satisfied with probability $1 - \prod_j (1-u_j)$.
- We claim that the above is minimized when $u_j = v_i/k$ for each j . (Take the proof as a reading exercise).
- So, the probability of interest is $1 - (1 - v_i/k)^k$.
- We now claim that the function $f(r) = 1 - (1 - r/k)^k$ is at least $1 - (1-1/k)^k \cdot r$ for all r in $[0,1]$.
 - Take the proof of the above also as a reading exercise. You need to show that the function is concave.

LP Rounding

- A clause C_i with k literals is satisfied with probability at least $1 - (1-1/k)^k \cdot v_i$.
- So, the probability of interest is $1 - (1 - v_i/k)^k$.
- We now claim that the function $f(r) = 1 - (1 - r/k)^k$ is at least $1 - (1-1/k)^k \cdot r$ for all r in $[0,1]$.
- By the above, we conclude that C_i is satisfied with probability at least $(1-1/k)^k \cdot v_i$.
- Now, use linearity of expectations (over clauses) to show that the expected number of satisfied clauses is at least $\sum_i (1 - (1-1/k)^k) \cdot v_i \geq (1 - 1/e) \sum_i v_i \geq (1-1/e) m^*(I)$.

LP Rounding

- There are two algorithms with performance guarantees as shown:

k	$1 - (1/2^k)$	$1 - (1 - 1/k)^k$
1	0.5	1.0
2	0.75	0.75
3	0.875	0.703
4	0.938	0.684
5	0.969	0.672

- As k , the length of the clause, increases, the first algorithm does better, and as k is small, the latter is better.

LP Rounding

- The next steps:
- We have two algorithms. One guarantees an approximation ratio of at least $1/2$.
- The other guarantees an approximation ratio of at least $1 - (1 - 1/k)^k$.
- We can use both to achieve an approximation ratio of at least $3/4$.
- Idea 1 : Run both algorithms and pick the best solution.
- Idea 2 : Toss a fair coin and pick an algorithm based on the outcome.....
- In both cases, we can show that the expected performance ratio will be at least $3/4$.

LP Rounding

- We have two algorithms. One guarantees an approximation ratio of at least $1/2$. The other guarantees an approximation ratio of at least $1 - (1 - 1/k)^k$.
- We can use both to achieve an approximation ratio of at least $3/4$.
- Idea 2 : Toss a fair coin and pick an algorithm based on the outcome.....
- In this case, the expected number of clauses satisfied is $(n_1 + n_2)/2$.
 - Here, n_1 and n_2 refer to the expected number of clauses satisfied by the first and the second algorithms.

LP Rounding

- We have two algorithms. One guarantees an approximation ratio of at least $1/2$. The other guarantees an approximation ratio of at least $1 - (1 - 1/k)^k$.
- Idea 2 : Toss a fair coin and pick an algorithm based on the outcome.....
- In this case, the expected number of clauses satisfied is $(n_1 + n_2)/2$.
 - Here, n_1 and n_2 refer to the expected number of clauses satisfied by the first and the second algorithms.
- We will show that $(n_1 + n_2)/2 \geq \frac{3}{4} \sum_i v_i$

LP Rounding

- In this case, the expected number of clauses satisfied is $(n_1 + n_2)/2$.
 - Here, n_1 and n_2 refer to the expected number of clauses satisfied by the first and the second algorithms.
- We will show that $(n_1 + n_2)/2 \geq \frac{3}{4} \sum_j v_j$
- Focus on clauses that have k literals. Then,
- $n_1 = \sum_k \sum_{C_j \text{ has } k \text{ literals}} (1 - 2^{-k}) \geq \sum_k \sum_{C_j \text{ has } k \text{ literals}} (1 - 2^{-k}) \cdot v_j$.
- We know that $n_2 \geq \sum_k \sum_{C_j \text{ has } k \text{ literals}} \beta_k \cdot v_j$ where $\beta_k = 1 - (1 - 1/k)^k$

LP Rounding

- We will show that $(n_1 + n_2)/2 \geq \frac{3}{4} \sum_j v_j$
- Focus on clauses that have k literals. Then,
- $n_1 = \sum_k \sum_{C_j \text{ has } k \text{ literals}} (1 - 2^{-k}) \geq \sum_k \sum_{C_j \text{ has } k \text{ literals}} (1 - 2^{-k}) \cdot v_j$.
- We know that $n_2 \geq \sum_k \sum_{C_j \text{ has } k \text{ literals}} \beta_k \cdot v_j$ where $\beta_k = 1 - (1 - 1/k)^k$
- So, $(n_1 + n_2)/2 \geq \sum_k \sum_{C_j \text{ has } k \text{ literals}} \frac{1}{2} [(1 - 2^{-k}) + \beta_k] \cdot v_j$.
- We prove that $(1 - 2^{-k}) + \beta_k$ is at least $3/2$.
- Ending the proof of the required claim.

Approximate Counting

- We now focus on a new topic called approximate counting.
- The idea here is to see how many objects satisfy a given condition.
- Examples:
 - Count the number of spanning trees of a given graph.
 - Count the number of matchings of a given graph.
 - Count the number of truth assignments that satisfy a formula in DNF.
- As can be seen, the first example has a polynomial time algorithm while the other two do not.

Approximate Counting

- We start with studying how to count the number of truth assignments that satisfy a formula in DNF.
- A formula in DNF is of the form $C_1 \vee C_2 \vee \dots \vee C_m$ where each clause C_i is of the form $(L_1 \wedge L_2 \wedge \dots \wedge L_j)$.
- The L 's are called literals and each literal is either a variable X_k or its negation X_k' .
- The problem has applications in network design and reliability.

Approximate Counting

- We start with studying how to count the number of truth assignments that satisfy a formula in DNF.
- A formula F in DNF is of the form $C_1 \vee C_2 \vee \dots \vee C_m$ where each clause C_i is of the form $(L_1 \wedge L_2 \wedge \dots \wedge L_{r_i})$.
- The L 's are called literals and each literal is either a variable X_k or its negation X_k' .
- A truth assignment is said to satisfy F if some clause in F is true under the assignment.
- We are interested in $\#F$ – the number of distinct satisfying assignments of F .
- While no deterministic algorithms are known, we focus on randomized algorithms to get as close to the answer as possible.

Approximate Counting

- We abstract the problem as follows.
- Let U be a finite set and $f : U \rightarrow \{0, 1\}$ be a Boolean function.
- Define $G = \{ u \mid f(u) = 1 \}$.
- We assume that given a u in U , $f(u)$ is easy to compute.
- We also assume that it is possible to sample uniformly at random from U .
- We want to estimate the size of G .

Approximate Counting

- The method we study is called the Monte Carlo method.
- Choose N independent samples u from U , say u_1, u_2, \dots, u_N .
- Evaluate f at each of these samples and count the number of instances that f evaluates to 1.
- Use this count to estimate G .
- Details follow.

Approximate Counting

- The method we study is called the Monte Carlo method.
- Choose N independent samples u from U , say u_1, u_2, \dots, u_N .
- Evaluate f at each of these samples and count the number of instances that f evaluates to 1.
- Use this count to estimate G .
- Let Y_i be a random variable that indicates whether u_i in G .
- Define another random variable $Z = |U| \cdot \sum_i Y_i / N$.
- Verify that $E[Z] = |G|$.

Approximate Counting

- The method we study is called the Monte Carlo method.
- Choose N independent samples u from U , say u_1, u_2, \dots, u_N .
- Evaluate f at each of these samples and count the number of instances that f evaluates to 1.
- Use this count to estimate G .
- Let Y_i be a random variable that indicates whether u_i is in G .
- Define another random variable $Z = |U| \cdot \sum_i Y_i / N$.
- Verify that $E[Z] = |G|$.
- $E[Z] = (|U|/N) \cdot \sum_i E[Y_i] = (|U|/N) \cdot \sum_i |G|/|U| = |G|$.

Approximate Counting

- The method we study is called the Monte Carlo method.
- Choose N independent samples u from U , say u_1, u_2, \dots, u_N .
- Evaluate f at each of these samples and count the number of instances that f evaluates to 1.
- Use this count to estimate G .
- Let Y_i be a random variable that indicates whether u_i is in G .
- Define another random variable $Z = |U| \cdot \sum_i Y_i / N$.
- Verify that $E[Z] = |G|$.
- $E[Z] = (|U|/N) \cdot \sum_i E[Y_i] = (|U|/N) \cdot \sum_i |G|/|U| = |G|$.
- We hope that since Z has an expectation that is useful for us, Z can be used as an estimator for $|G|$.

Approximate Counting

- Define another random variable $Z = |U| \cdot \sum_i Y_i / N$.
- Verify that $E[Z] = |G|$.
- $E[Z] = (|U|/N) \cdot \sum_i E[Y_i] = (|U|/N) \cdot \sum_i |G|/|U| = |G|$.
- We hope that since Z has an expectation that is useful for us, Z can be used as an estimator for $|G|$.
- In other words, we hope that Z has a value close to $|G|$.
- Indeed, an application of Chernoff bounds again.

Approximate Counting

- Let $r = |G|/|U|$. Then, with probability at least $1 - \delta$, the Monte Carlo estimation produces a result that is within ε of the actual value.
- Recall that each Y_i is a Bernoulli random variable and the Y_i 's are iid.
- Define $Y = \sum_i Y_i$ and realize that $Z = |U|Y/N$.
- $E[Y] = Nr$.
- Using Chernoff bounds, $\Pr((1-\varepsilon)|G| \leq Z \leq (1+\varepsilon)|G|)$
- $= \Pr((1-\varepsilon)|G| \leq |U|Y/N \leq (1+\varepsilon)|G|)$
- $= \Pr((1-\varepsilon)Nr \leq Y \leq (1+\varepsilon)Nr)$
- $\geq 1 - 2\exp\{-\varepsilon^2 Nr/4\}$.

Approximate Counting

- Let $r = |G|/|U|$. Then, with probability at least $1 - \delta$, the Monte Carlo estimation produces a result that is within ε of the actual value with $N \geq 4 \ln(2/\delta)/\varepsilon^2 r$.
- Define $Y = \sum_i Y_i$ and realize that $Z = |U|Y/N$. $E[Y] = Nr$.
- Using Chernoff bounds, $\Pr((1-\varepsilon)|G| \leq Z \leq (1+\varepsilon)|G|)$
- $= \Pr((1-\varepsilon)|G| \leq |U|Y/N \leq (1+\varepsilon)|G|)$
- $= \Pr((1-\varepsilon)Nr \leq Y \leq (1+\varepsilon)Nr)$
- $\geq 1 - 2\exp\{-\varepsilon^2 Nr/4\}$.
- With the choice of N , $\exp\{-\varepsilon^2 Nr/4\} = \exp\{-\ln(2/\delta)\}$.
- So, the required probability is at least $1 - \delta$.

Approximate Counting

- How good is our scheme?
- Notice that the scheme requires N samples to be drawn.
- But $N = O(1/r)$ where $r = |G|/|U|$.
- If $|G|$ is small, then N is large.
- N is not only large, but is as large as $1/r$ implying that N can get close to being exponentially large.

Approximate Counting

- How good is our scheme?
- Notice that the scheme requires N samples to be drawn.
- But $N = O(1/r)$ where $r = |G|/|U|$.
- If $|G|$ is small, then N is large.
- N is not only large, but is as large as $1/r$ implying that N can get close to being exponentially large.
- The trouble is not with the Chernoff bound but with u ar sampling.
- For cases with r too small, fetching useful samples is difficulty. So, needs more trails.

Importance Sampling

- Do not rely on uniform sampling but use skewed sampling.
- Skew towards useful samples.
- Details follow.

Importance Sampling

- One aspect of importance sampling is to increase the value of $r = |G|/|U|$.
- With a large r , we will avoid the danger of $N = O(1/r)$ getting large when r is small.
- The approach we study is quite generic.
- Let V be a finite universe.
- We are given m subsets H_1, H_2, \dots, H_m of V such that
 - 1) For all i , $|H_i|$ can be obtained in polynomial time.
 - 2) It is possible to sample uniformly at random from any H_i .
 - 3) For all v in V , it is possible to determine if v in H_i in polynomial time.

Importance Sampling

- One aspect of importance sampling is to increase the value of $r = |G|/|U|$.
- With a large r , we will avoid the danger of $N = O(1/r)$ getting large when r is small.
- The approach we study is quite generic.
- Let V be a finite universe.
- We are given m subsets H_1, H_2, \dots, H_m of V such that
 - 1) For all i , $|H_i|$ can be obtained in polynomial time.
 - 2) It is possible to sample uniformly at random from any H_i .
 - 3) For all v in V , it is possible to determine if v in H_i in polynomial time.
- The goal is to estimate the size of $H = \bigcup_i H_i$.

Importance Sampling

- Let V be a finite universe.
- We are given m subsets H_1, H_2, \dots, H_m of V such that
 - 1) For all i , $|H_i|$ can be obtained in polynomial time.
 - 2) It is possible to sample uniformly at random from any H_i .
 - 3) For all v in V , it is possible to determine if v in H_i in polynomial time.
- The goal is to estimate the size of $H = \bigcup_i H_i$.
- One possibility is to use the inclusion-exclusion principle.
- But there are too many terms in the formula.

Importance Sampling

- Let V be a finite universe. We are given m subsets H_1, H_2, \dots, H_m of V such that:
 - 1) For all i , $|H_i|$ can be obtained in polynomial time.
 - 2) It is possible to sample uniformly at random from any H_i .
 - 3) For all v in V , it is possible to determine if v is in H_i in polynomial time.
- The goal is to estimate the size of $H = \bigcup_i H_i$.
- How is the DNF Counting problem similar?
- The universe V is the set of all possible assignments.
- Each H_i corresponds to the assignments that satisfy C_i .
- What is $|H_i|$ for a given i ?

Importance Sampling

- Let V be a finite universe. We are given m subsets H_1, H_2, \dots, H_m of V such that:
 - 1) For all i , $|H_i|$ can be obtained in polynomial time.
 - 2) It is possible to sample u from any H_i .
 - 3) For all v in V , it is possible to determine if v in H_i in polynomial time.
- The goal is to estimate the size of $H = \bigcup_i H_i$.
- How is the DNF Counting problem similar?
- The universe V is the all possible assignments.
- Each H_i corresponds to the assignments that satisfy C_i .
- What is $|H_i|$ for a given i ? Exactly 2^{n-r_i} .
- Easy to sample too. Fix the truth values of r_i variables, others can be set u .

Importance Sampling

- Let V be a finite universe. We are given m subsets H_1, H_2, \dots, H_m of V such that:
 - 1) For all i , $|H_i|$ can be obtained in polynomial time.
 - 2) It is possible to sample u from any H_i .
 - 3) For all v in V , it is possible to determine if v is in H_i in polynomial time.
- The goal is to estimate the size of $H = \bigcup_i H_i$.
- How is the DNF Counting problem similar?
- The universe V is the all possible assignments.
- Each H_i corresponds to the assignments that satisfy C_i .
- What is $|H_i|$ for a given i ? Exactly 2^{n-r_i} .
- Easy to sample too. Fix the truth values of r_i variables, others can be set u .
- Can also check if a given v in V is in H_i .

Importance Sampling

- Let V be a finite universe. We are given m subsets H_1, H_2, \dots, H_m of V such that:
 - 1) For all i , $|H_i|$ can be obtained in polynomial time.
 - 2) It is possible to sample u from any H_i .
 - 3) For all v in V , it is possible to determine if v is in H_i in polynomial time.
- The goal is to estimate the size of $H = \bigcup_i H_i$.
- How is the DNF Counting problem similar?
- The universe V is the all possible assignments.
- Each H_i corresponds to the assignments that satisfy C_i .
- What is $|H_i|$ for a given i ? Exactly 2^{n-r_i} .
- Easy to sample too. Fix the truth values of r_i variables, others can be set u .
- Can also check if a given v in V is in H_i .
- And, the DNF counting problem is to estimate $|H|$.

Importance Sampling

- Define a multiset U as the union of H_i 's where the multiset keeps multiple copies of each element.
- Viewed differently, U can be seen as consisting of tuples (v, i) such that v is in H_i .
- Hence, $U = \{ (v, i) \mid v \text{ is in } H_i \}$.
- Now, $|U| = \sum_i |H_i| \geq |H|$.
- We also define for every v in V , the coverage set of v as
- $\text{cov}(v) = \{ (v, i) \mid (v, i) \text{ in } U \}$
- Notice that for a given v , the size of $\text{cov}(v)$ is the number of H_i 's that contain v_i .

Importance Sampling

- Viewed differently, U can be seen as consisting of tuples (v, i) such that v is in H_i .
- Hence, $U = \{ (v, i) \mid v \text{ is in } H_i \}$. Now, $|U| = \sum_i |H_i| \geq |H|$.
- We also define for every v in V , the coverage set of v as
- $\text{cov}(v) = \{ (v, i) \mid (v, i) \text{ in } U \}$. Notice that for a given v , the size of $\text{cov}(v)$ is the number of H_i 's that contain v_i .
- The following observations hold wrt $\text{cov}()$
 - 1) The number of coverage sets is exactly $|H|$.
 - 2) The coverage sets partition U , i.e. $U = \bigcup_{v \text{ in } H} \text{cov}(v)$.
 - 3) As a result of 2, $|U| = \sum_{v \text{ in } H} |\text{cov}(v)|$.
 - 4) For all v in H , $|\text{cov}(v)|$ is at most m .

Importance Sampling

- Based on the above definitions, for each coverage set, we define a canonical element as follows.
- Define a function $f: U \rightarrow \{0,1\}$ as $f((v,i)) = 1$ if $i = \min\{j \mid v \text{ in } H_j\}$ and 0 otherwise.
- Define $G = \{ (v,j) \text{ in } U \mid f((v,j)) = 1 \}$.
- Notice that $|G| = |H|$.

Importance Sampling

- Based on the above definitions, for each coverage set, we define a canonical element as follows.
- Define a function $f: U \rightarrow \{0,1\}$ as $f((v,i)) = 1$ if $i = \min\{j \mid v \text{ in } H_j\}$ and 0 otherwise.
- Define $G = \{ (v,j) \text{ in } U \mid f((v,i)) = 1 \}$.
- Notice that $|G| = |H|$.
- Further, $r = |G|/|U| \geq 1/m$ now. As,
 - $|U| = \sum_{v \text{ in } H} |\text{cov}(v)| \leq \sum_{v \text{ in } H} m \leq m |H| = m |G|$.
- Now, use the Monte-Carlo technique of sampling on U .
- Same calculations as earlier hold now. Except that the time taken is $O(1/r) = O(\text{polynomial})$.
- Only have to show how to sample u from U .

Importance Sampling

- Sampling uar from U is done in two steps as follows:
 - 1) Choose an i between 1 and m such that $\Pr[i \text{ is chosen}] = |H_i|/|U|$.
 - 2) Once an i is chosen, choose a v in H_i u.a.r.
- Can claim that the pair (v,i) is now uar over U .

Importance Sampling

- Why did this sampling procedure work?
- Cannot sample u from V and quickly get a good approximation.
- However, using U in the importance sampling experiment is equivalent to sampling a random satisfying truth assignment from a randomly chosen clause.
- The probability we choose a particular assignment now depends on the number of clauses this assignment satisfies.
- However, since G is defined on the basis of canonical tuples, only one choice of truth assignment and the correct clause makes it count correctly towards the size of G .

Some Definitions

- With respect to approximate counting algorithms, one often uses the following notion.
- A polynomial randomized approximation scheme (PRAS) for a counting problem Π is a randomized algorithm A running in polynomial time (in n) such that on any instance I of P and a real number $\varepsilon > 0$, produces an output $A(I)$ such that
$$\Pr[(1 - \varepsilon) \#I \leq A(I) \leq (1 + \varepsilon) \#I] \geq 3/4.$$
- A PRAS scheme for a problem Π is called a fully polynomial randomized approximation scheme (FPRAS) if in addition the runtime of the algorithm A is bounded by a polynomial in n and $1/\varepsilon$.

Some Definitions

- A polynomial randomized approximation scheme (PRAS) for a counting problem Π is a randomized algorithm A running in polynomial time (in n) such that on any instance I of P and a real number $\varepsilon > 0$, produces an output $A(I)$ such that
$$\Pr[(1 - \varepsilon) \#I \leq A(I) \leq (1 + \varepsilon) \#I] \geq 3/4.$$
- A PRAS scheme for a problem Π is called a fully polynomial randomized approximation scheme (FPRAS) if in addition the runtime of the algorithm A is bounded by a polynomial in n and $1/\varepsilon$.
- If the probability above can be made arbitrarily small to a given δ , while the runtime is $\text{poly}(n, 1/\varepsilon, \delta)$, then such a scheme is called as an (ε, δ) -FPRAS.

Some Definitions

- A problem is in the class $\#P$ if the corresponding decision problem is in NP.
- More formally, a problem is in $\#P$ if there is a non-deterministic polynomial time Turing machine that for any instance I has a number of accepting computations that is exactly equal to the number of distinct solution to instance I .
- Further, a problem Π is said to be $\#P$ -complete if for any problem Π' in $\#P$, Π' can be reduced to Π in polynomial time.

